



# HashiCorp Terraform Associate Lecture Slides

These lecture slides are provided for personal and non-commercial use only

Please do not redistribute or upload these lecture slides elsewhere.

Good luck on your exam!

# What is Infrastructure as Code?

## The Problem with Manual Configuration

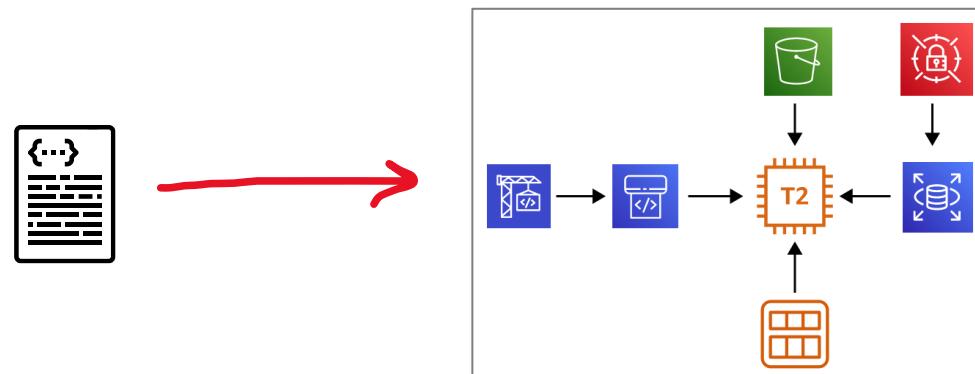
Manually configuring your cloud infrastructure allows you easily start using new service offerings to quickly prototype architectures however it comes with many downsides:

- Its easy to mis-configure a service though human error
- Its hard to manage the expected state of configuration for compliance
- Its hard to transfer configuration knowledge to other team members

## Infrastructure as Code (IaC)

You write a configuration script to **automate** **creating, updating or destroying** cloud infrastructure.

- IaC is a **blueprint** of your infrastructure.
- IaC allows you to easily **share, version or inventory** your cloud infrastructure.



# Popular Infrastructure as Code tools (IaC)

## Declarative

- What you see is what you get. *Explicit*
- More verbose, but zero chance of mis-configuration
- Uses scripting languages eg. JSON, YAML, XML



### ARM Templates

Supports only Azure



### Azure Blueprints

Supports only Azure

Manages relationship between services



### CloudFormation

Only for AWS



### Cloud Deployment Manager

Supports on Google Cloud



### Terraform

Supports many cloud service providers (CSPs) and cloud services.

## Imperative

- You say what you want, and the rest is filled in. *Implicit*
- Less verbose, you could end up with misconfiguration
- Does more than Declarative
- Uses programming languages eg. Python, Ruby, JavaScript



### AWS Cloud Development Kit (CDK)

Supports only AWS

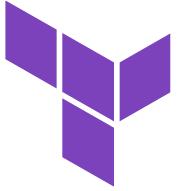
Many built-in templates for opinionated best practices



### Pulumi

Supports AWS, Azure, GCP, K8

# Declarative+



**Terraform** is declarative but the Terraform Language features **imperative-like functionality**.

| Declarative                                                                                        |                                                                                                                                                                                                                  | Imperative                                                                                |
|----------------------------------------------------------------------------------------------------|------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|-------------------------------------------------------------------------------------------|
| YAML, JSON, XML                                                                                    | HCL-ish (Terraform Language)                                                                                                                                                                                     | Ruby, Python, JavaScript ...                                                              |
| Limited or no support for imperative-like features.                                                | Supports:                                                                                                                                                                                                        | Imperative features is the utility of the entire feature set of the programming language. |
| In some cases you can add behaviour by extending the base language.<br>E.g. CloudFormation Macros. | <ul style="list-style-type: none"><li>• Loops (For Each)</li><li>• Dynamic Blocks</li><li>• Locals</li><li>• Complex Data Structure<ul style="list-style-type: none"><li>• Maps, Collections</li></ul></li></ul> |                                                                                           |

# Infrastructure Lifecycle

## What is Infrastructure Lifecycle?

a number **of clearly defined and distinct work phases** which are used by DevOps Engineers to **plan, design, build, test, deliver, maintain and retire** cloud infrastructure.

## What is Day 0, Day 1 and Day 2?

Day 0-2 is a simplified way to describe phases of an infrastructure lifecycle

- Day 0 — Plan and Design
- Day 1 — Develop and Iterate
- Day 2 — Go live and maintain

Days do not literally mean a **24 hour days** and is just a broad way of defining where a Infrastructure project would be.

# Infrastructure Lifecycle

How does IaC enhance the Infrastructure Lifecycle?

## Reliability

IaC makes changes **idempotent**, consistent, repeatable, and predictable.

### **Idempotent**

No matter how many times you run IaC, you will always end up with the same state that is expected

## Manageability

- enable mutation via code
- revised, with minimal changes

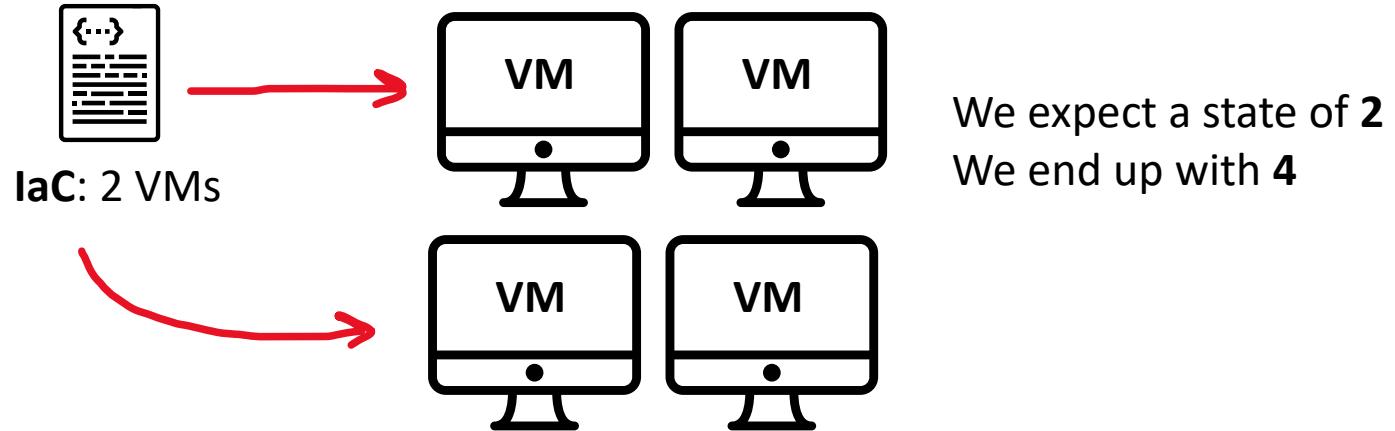
## Sensibility

avoid financial and reputational losses to even loss of life when considering government and military dependencies on infrastructure.

# Non-Idempotent vs Idempotent

## Non-Idempotent

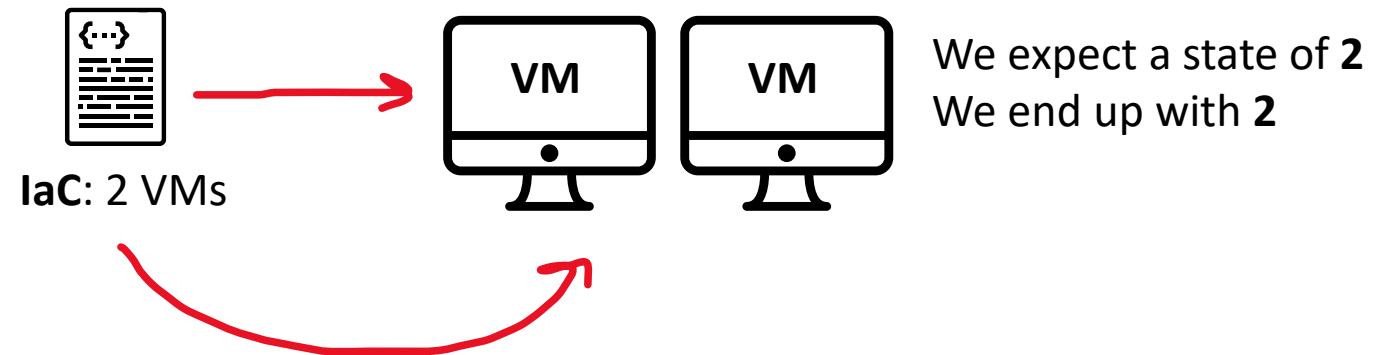
When I deploy my IaC config file it will **provision** and **launch** 2 virtual machines



When I update my IaC and deploy again, I will end up with 2 new VMs with a total of 4 VMs

## Idempotent

When I deploy my IaC config file it will **provision** and **launch** 2 virtual machines



When I update my IaC and deploy again, it will update the VMs if changed by **modifying or deleting and creating new VMs**

# Provisioning vs Deployment vs Orchestration

## Provisioning

To prepare a server with systems, data and software, and make it ready for network operation. Using Configuration Management tools like Puppet, Ansible, Chef, Bash scripts, PowerShell or Cloud-Init you can provision a server.

**When you launch a cloud service and configure it you are “provisioning”**

## Deployment

Deployment is the act of delivering a version of your application to run a provisioned server. Deployment could be performed via AWS CodePipeline, Harness, Jenkins, Github Actions, CircleCI

## Orchestration

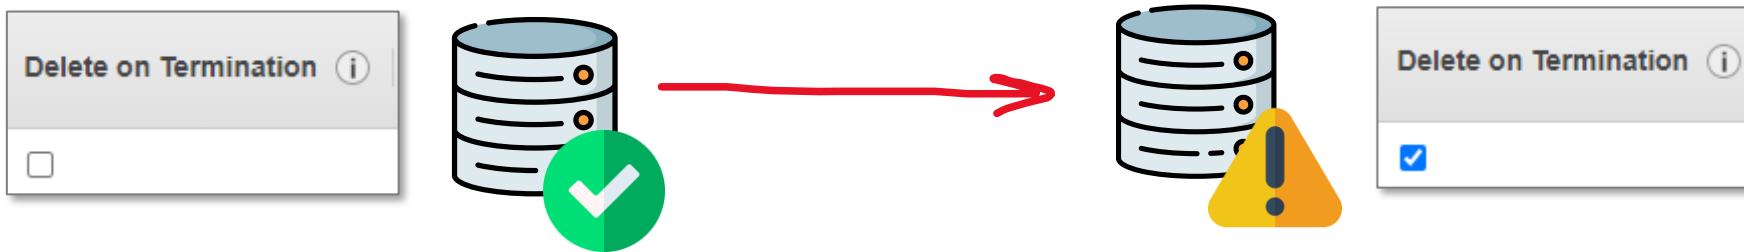
Orchestration is the act of coordinating multiple systems or services. Orchestration is a common term when working with microservices, Containers and Kubernetes. Orchestration could be Kubernetes, Salt, Fabric

# Configuration Drift

**Configuration Drift** is when provisioned infrastructure has  
**an unexpected configuration change** due to:

- team members manually adjusting configuration options
- malicious actors
- side affects from APIs, SDK or CLIs.

eg, a junior developer turns on Delete on Termination for the production database.



Configuration Drift going unnoticed could be loss or breach of cloud services and residing data or result in interruption of services or unexpected downtime.

# Configuration Drift

## How to detect configuration drift?

- A compliance tool that can detect misconfiguration eg. AWS Config, Azure Policies, \*GCP Security Health Analytics
- Built-in support for drift detection eg. AWS CloudFormation Drift Detection
- Storing the expected state eg. Terraform state files

## How to correct configuration drift?

- A compliance tool that can remediate (correct) misconfiguration e.g. AWS Config
- Terraform refresh and plan commands
- Manually correcting the configuration (not recommended)
- Tearing down and setting up the infrastructure again

## How to prevent configuration drift?

- **Immutable infrastructure**, always create and destroy, never reuse, Blue, Green deployment strategy.
  - Servers are never modified after they are deployed
    - Baking AMI images or containers via AWS Image Builder or HashiCorp Packer, or a build server eg. GCP Cloud Run
- Using GitOps to version control our IaC, and peer review every single via Pull Requests change to infrastructure

# Mutable vs Immutable Infrastructure

## Mutable Infrastructure



A Virtual Machine (VM) is deployed and then a Configuration Management tool like Ansible, Puppet, Chef, Salt or Cloud-Init is used to configure the state of the server



## Immutable Infrastructure



A VM is launched and provisioned, and then it is turned into a Virtual Image, stored in image repository, that image is used to deployed VM instances



# Immutable Infrastructure Guarantee

Terraform encourages you towards an Immutable Infrastructure architect so you get the following guarantees.

**Cloud Resource Failure** – What if an EC2 instance fails a status check?

**Application Failure** – What if your post installation script fails due to change in package?

**Time to Deploy** - What if I need to deploy in a hurry?

## Worst Case Scenario

- Accidental Deletion
- Compromised by malicious actor
- Need to Change Regions (region outage)

## No Guarantee of 1-to-1

Every time Cloud-Init runs post deploy there is no guarantee its one-to-one with your other VMs.



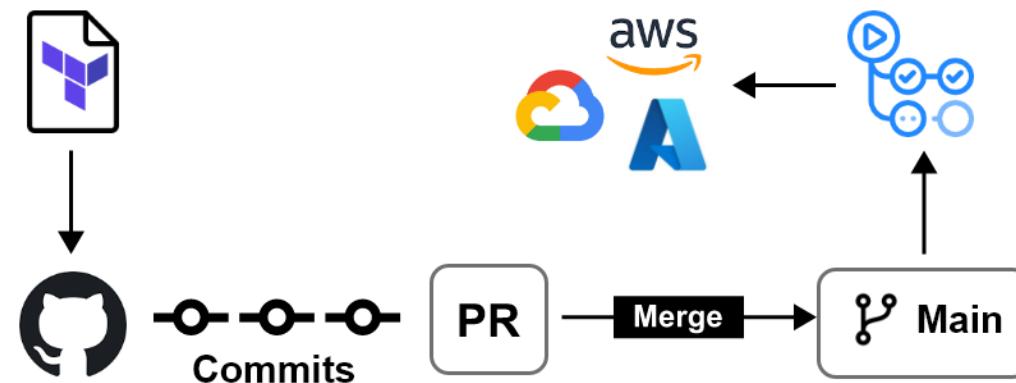
Packer

## Golden Images

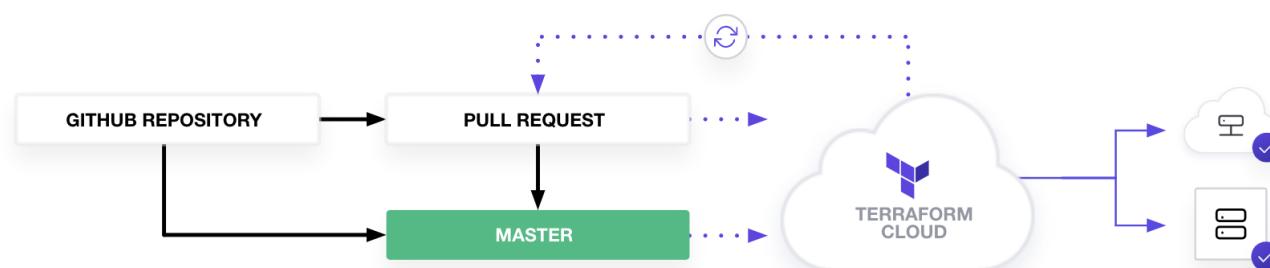
- Guarantee of 1-to-1 with your fleet
- Increased assurance of consistency, security
- Speeds up your deployments

# What is GitOps?

GitOps is when you take Infrastructure as Code (IaC) and you use a git repository to introduce a formal process to review and accept changes to infrastructure code, once that code is accepted, it automatically triggers a deploy



TERRAFORM CLOUD AND GITHUB ACTIONS WORKFLOW



# HashiCorp



HashiCorp is a company specializing in managed **open-source tools** used to support the **development and deployment of large-scale service-oriented software installations**

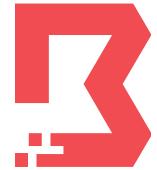
## What is HashiCorp Cloud Platform (HCP)?

HCP is a unified platform to access Hashicorp various products.

HCP services are **cloud agnostic**

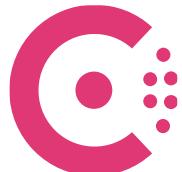
- support for the main cloud service providers (CSPs)
  - eg. AWS, GCP and Azure
- highly suited for **multi-cloud** workloads

# HashiCorp Products



## Boundary

secure remote access to systems based on trusted identity.



## Consul

service discovery platform. provides a full-featured service mesh for secure service segmentation across any cloud or runtime environment, and distributed key-value storage for application configuration



## Nomad

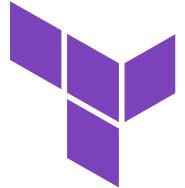
scheduling and deployment of tasks across worker nodes in a cluster



## Packer

tool for building virtual-machine images for later deployment.

# HashiCorp Products



## Terraform

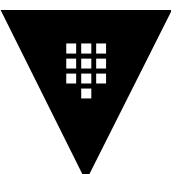
infrastructure as code software which enables provisioning and adapting virtual infrastructure across all major cloud provider

**Terraform Cloud** – a place to storage and manage IaC in the cloud or with teams



## Vagrant

building and maintenance of reproducible software-development environments via virtualization technology



## Vault

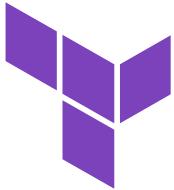
secrets management, identity-based access, encrypting application data and auditing of secrets for applications, systems, and users



## Waypoint

modern workflow to build, deploy, and release across platforms

# What is Terraform?



Terraform is an **open-source** and **cloud-agnostic** Infrastructure as Code (IaC) tool.

Terraform uses **declarative** configuration files.

The configuration files are written in  
**HashiCorp Configuration Language (HCL)**.

## Notable features of Terraform:

- Installable modules
- Plan and predict changes
- Dependency graphing
- State management
- Provision infrastructure in familiar languages
  - via AWS CDK
- Terraform Registry with 1000+ providers



```
resource "aws_instance" "iac_in_action" {  
    ami           = var.ami_id  
    instance_type = var.instance_type  
    availability_zone = var.availability_zone  
  
    // dynamically retrieve SSH Key Name  
    key_name = aws_key_pair.iac_in_action.key_name  
  
    // dynamically set Security Group ID (firewall)  
    vpc_security_group_ids = [aws_security_group.iac_in_action.id]  
  
    tags = {  
        Name = "Terraform-managed EC2 Instance for IaC in Action"  
    }  
}
```

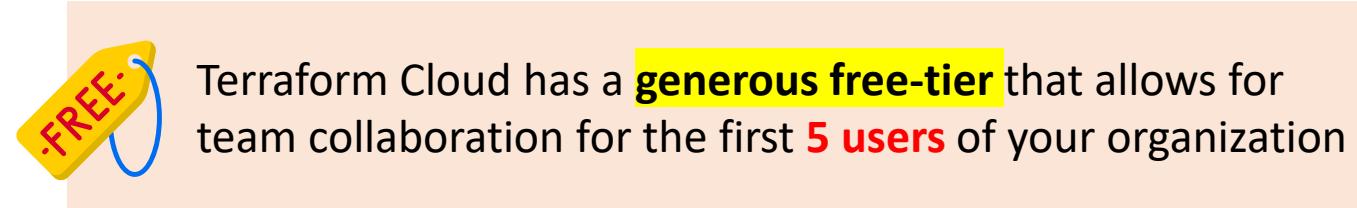
# What is Terraform Cloud?

The screenshot shows the Terraform Cloud interface for a workspace named "example-workspace". The top navigation bar includes links for ExamPro, Workspaces, Registry, Settings, and HCP. The main content area displays the workspace's status (0 resources, Terraform version 1.0.4, updated 18 days ago), execution mode (Remote), and auto-apply settings (Off). A "Metrics" section provides performance data over the last 20 runs. The "Latest Run" section highlights a successful refresh-only run triggered by a user named andrewbrown via the UI, completed 18 days ago. The "Resources" and "Outputs" sections show 0 items each. The "Tags" section notes that no tags have been added to the workspace.

Terraform Cloud is a Software as Service (SaaS) offering for:

- Remote state storage
- Version Control integrations
- Flexible workflows
- Collaborate on Infrastructure changes in a single **unified web portal**.

[www.terraform.io/cloud](https://www.terraform.io/cloud)

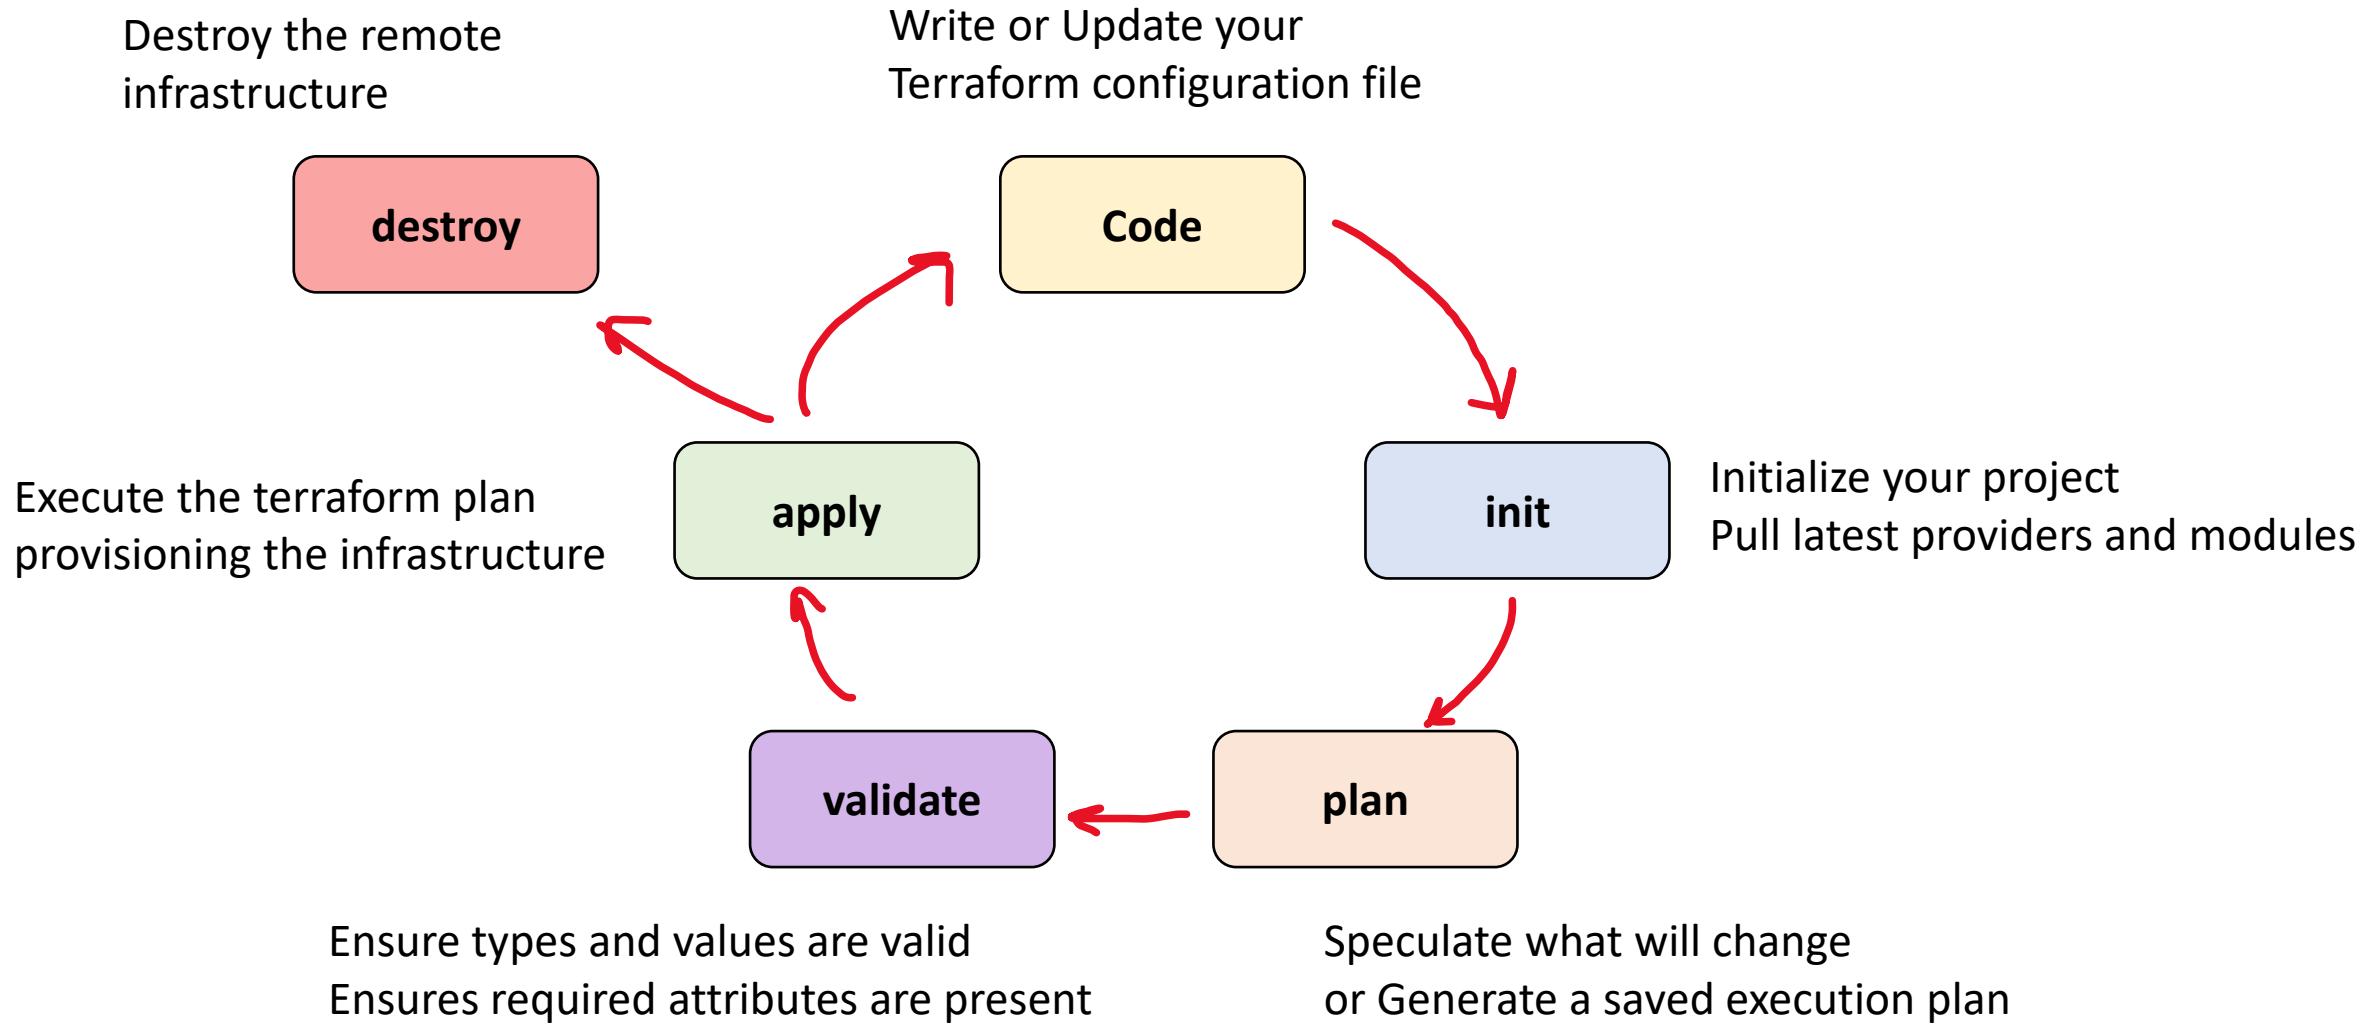


In majority of cases you should be using Terraform Cloud.

The only case where you may not want to use it to manage your state file is your company has many regulatory requirements along with a long procurement process so you could use Standard remote backend, Atlantis, or investigate Terraform Enterprise

The underlying software for Terraform Cloud and Terraform Enterprise is known as the “Terraform Platform”

# Terraform Lifecycle



# Terraform – Change Automation

## What is Change Management?

A standard approach to apply change, and resolving conflicts brought about by change. In the context of Infrastructure as Code (IaC), Change management is the procedure that will be followed when resources are modified and applied via configuration script.

## What is Change Automation?

a way of **automatically** creating a consistent, systematic, and predictable way of managing change request via controls and policies

Terraform uses Change Automation in the form of **Execution Plans** and **Resources graphs** to apply and review complex **changesets**

## What is a ChangeSet?

A collection of commits that represent changes made to a versioning repository. IaC uses ChangeSets so you can see what has changed by who over time.

Change Automation allows you to know exactly what Terraform will change and in what order, avoiding many possible human errors.

# Terraform - Execution Plans

An **Execution Plan** is a **manual review** of what will **add, change or destroy** before you apply changes eg. `terraform apply`

resources and configuration settings will be listed.

It will indicate will be added, changed or destroyed if this plan is approved:

Terraform used the selected providers to generate the following execution plan. Resource actions are indicated with the following symbols:

+ create

Terraform will perform the following actions:

```
# aws_instance.app_server will be created
+ resource "aws_instance" "app_server" {
    + ami                               = "ami-830c94e3"
    + arn                             = (known after apply)
    + associate_public_ip_address     = (known after apply)
    + availability_zone                = (known after apply)
    + cpu_core_count                  = (known after apply)
    + cpu_threads_per_core           = (known after apply)
    + disable_api_termination        = (known after apply)
    + ebs_optimized                  = (known after apply)
    + get_password_data              = false
    ...
}
```

Plan: 1 to add, 0 to change, 0 to destroy.

Do you want to perform these actions?

Terraform will perform the actions described above.

Only 'yes' will be accepted to approve.

Enter a value: **yes**

A user must approve changes by typing: **yes**

# Terraform – Visualizing Execution Plans

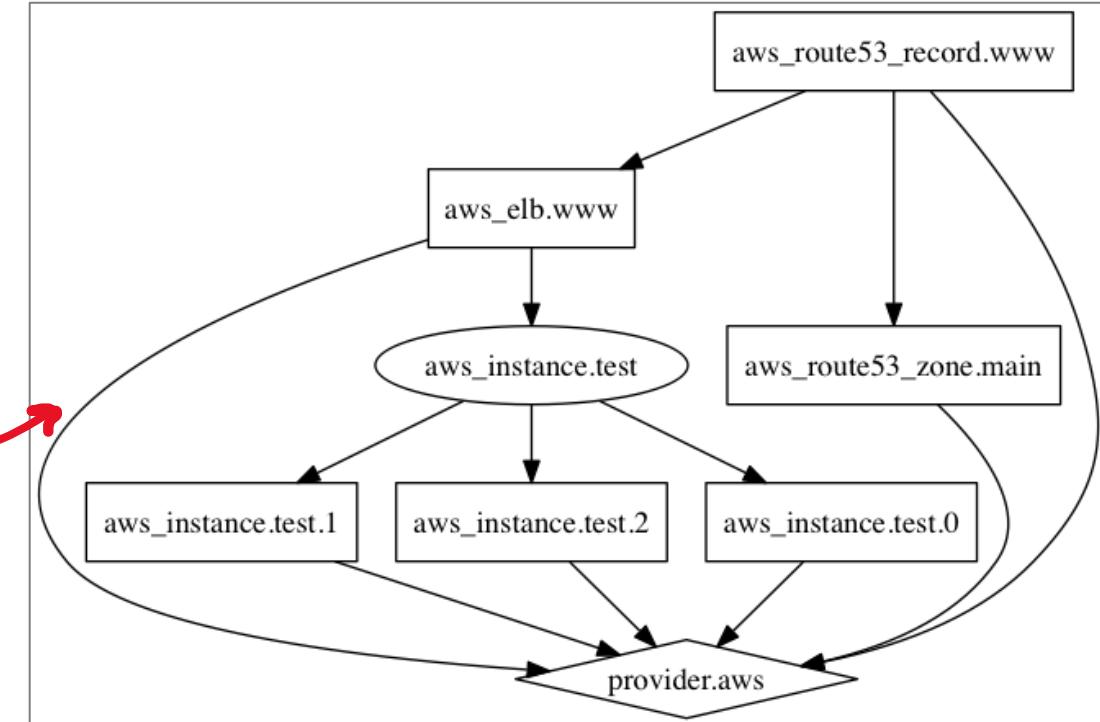
You can **visualize an execution plan** as a graph using the **terraform graph** command  
Terraform will output a GraphViz file (you'll need GraphViz installed to view the file)



## What is GraphViz?

open-source tools for drawing graphs specified in DOT language scripts having the file name extension "gv"

```
terraform graph | dot -Tsvg > graph.svg
```



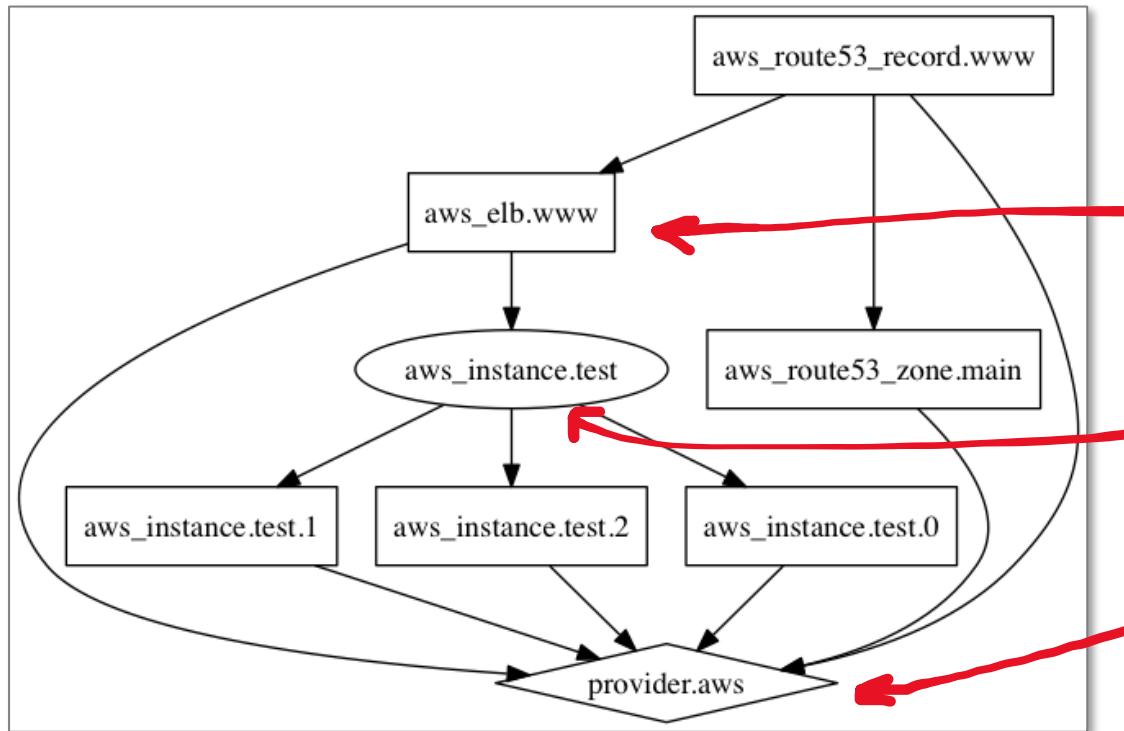
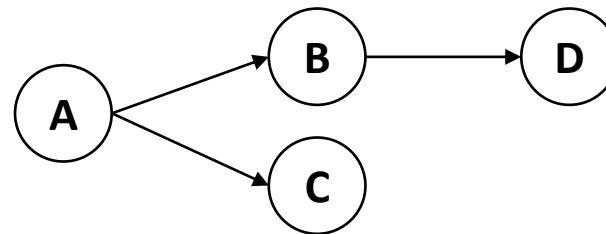
# Terraform – Resource Graph

Terraform builds a **dependency graph** from the Terraform configurations, and walks this graph to generate plans, refresh state, and more.

When you use **terraform graph**, this is a **visual presentation** of the dependency graph

## What is a dependency graph?

In mathematics is a directed graph representing dependencies of several objects towards each other



**Resource Node**  
Represents a single resource

**Resource Meta-Node**  
Represents a group of resources, but does not represent any action on its own

**Provider Configuration Node**  
Represents the time to fully configure a provider

# Terraform – Use Cases

## IaC for Exotic Providers

Terraform supports a variety of providers outside of GCP, AWS, Azure and sometimes is the only provider. Terraform is open-source and extendable so any API could be used to create IaC tooling for any kind of cloud platform or technology. E.g. Heroku, Spotify Playlists.

**Multi-Tier Applications** – Terraform by default makes it easy to divide large and complex applications into isolate configuration scripts (module). It has a complexity advantage over cloud-native IaC tools for its flexibility while retaining simplicity over Imperative tools.

**Disposable Environments** – Easily stand up an environment for a software demo or a temporary development environment

**Resource Schedulers** – Terraform is not just defined to infrastructure of cloud resource but can be used to dynamic schedule Docker containers, Hadoop, Spark and other software tools. You can provision your own scheduling grid.

**Multi-Cloud Deployment** – Terraform is cloud-agnostic and allows a single configuration to be used to manage multiple providers, and to even handle cross-cloud dependencies.

# Terraform Core and Terraform Plugins

Terraform is logically split into two main parts:

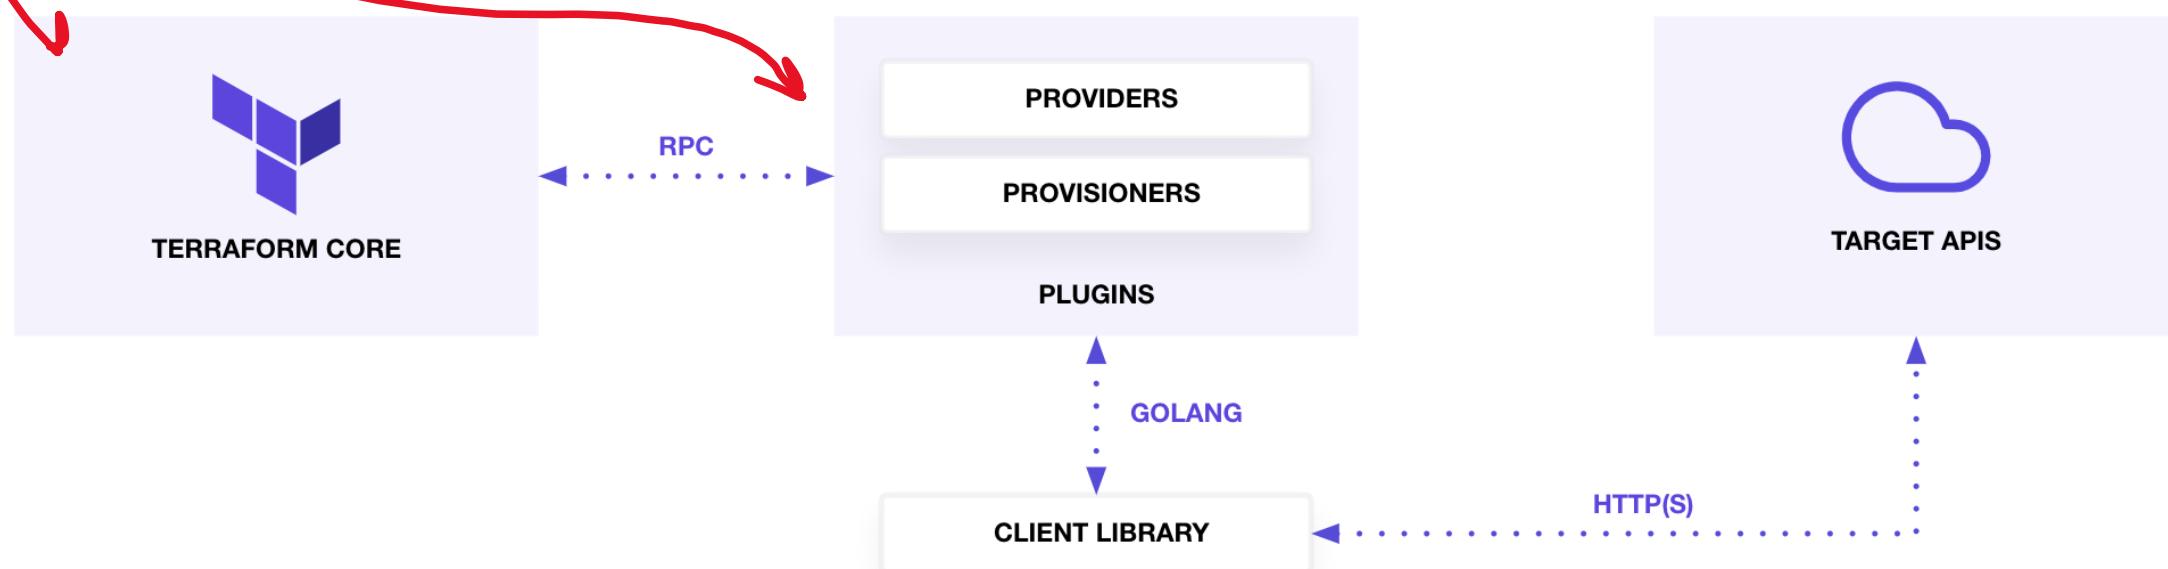
1. **Terraform Core**

- uses remote procedure calls (RPC) to communicate with Terraform Plugins

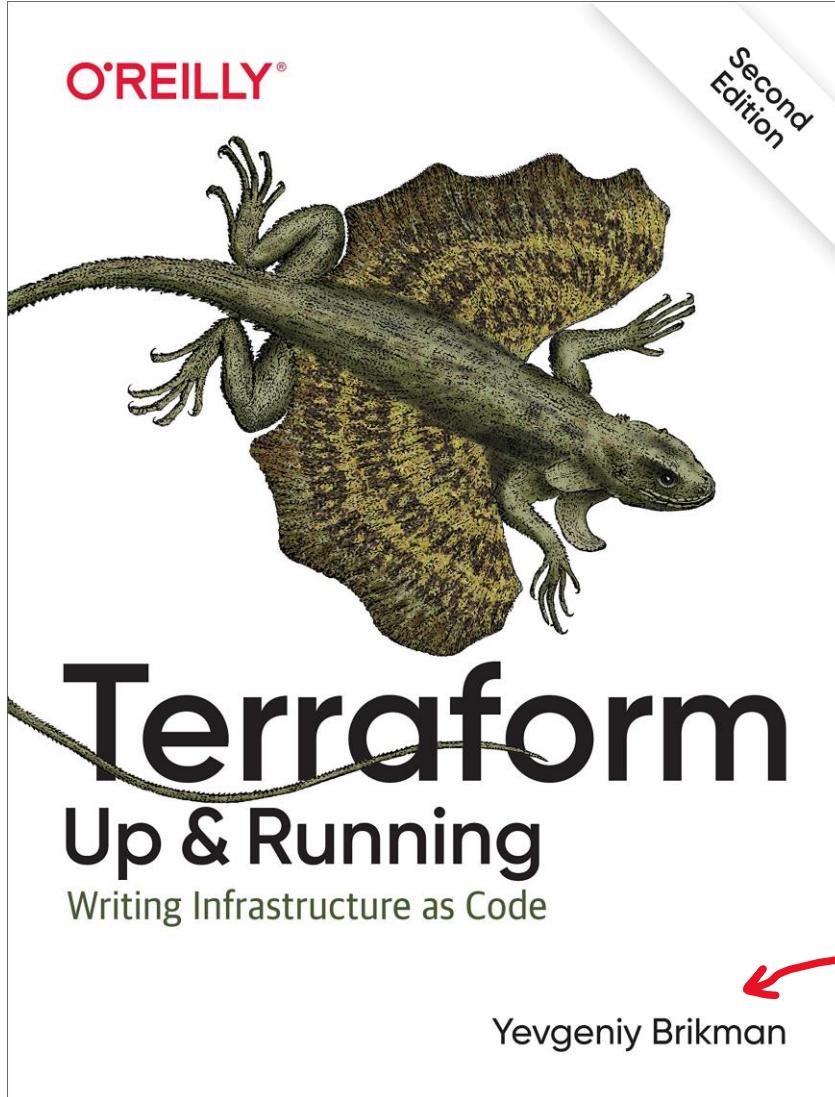
2. **Terraform Plugins**

- expose an implementation for a specific service, or provisioner

Terraform Core is a statically-compiled **binary** written in the Go programming language.



# Terraform Up and Running



**Terraform Up & Running** takes a deep dive into the internal workings of Terraform.

If you want to go beyond this course for things like:

- Testing your Terraform Code
- Zero-Downtime Deployment
- Common Terraform Gotchas
- Composition of Production Grade Terraform Code

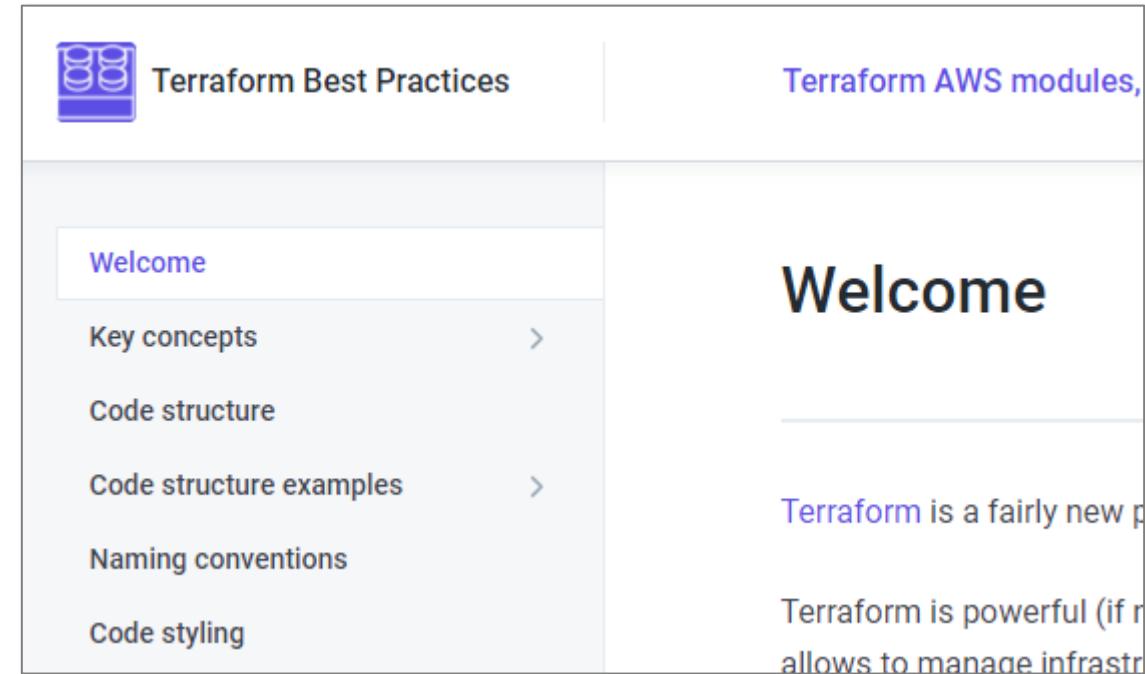
Yevgeniy ("Jim") Brikman is the co-founder of **Gruntwork**



# Terraform Best Practices

Terraform is an **online open-book**  
about **Terraform Best Practices**

[www.terraform-best-practices.com](http://www.terraform-best-practices.com)



The screenshot shows the homepage of the Terraform Best Practices website. At the top, there's a header with the title "Terraform Best Practices" and a "Terraform AWS modules" link. Below the header is a navigation menu with links to "Welcome", "Key concepts", "Code structure", "Code structure examples", "Naming conventions", and "Code styling". To the right of the menu, a portion of the "Welcome" page is visible, featuring the word "Welcome" in large letters and some descriptive text about Terraform.

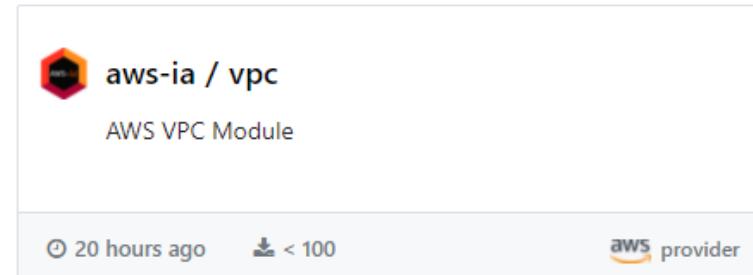
# AWS-IA vs Terraform AWS Modules



## AWS IA **Not Recommended**

Authored by AWS Integration and Automation team

- They don't follow best practices for Terraform module development
- Very limited in functionality
- Very few libraries
- Are new, do not have much use



## Terraform AWS Modules **Recommended**

Authored by Anton Babenko

- They follow best practices for Terraform Modules development
- Very flexible, covering many different use cases
- Many different libraries
- Have been around for **5 years**, heavily used (even previously by AWS)



↑  
**10M!**

# Atlantis



Atlantis is an open-source developer tool to  
**automate Terraform Pull Requests**

<https://www.runatlantis.io>

The screenshot shows a GitHub pull request interface. At the top, there's a comment from 'atlantisbot' with a profile picture of the Atlantis logo. The comment text is:

```
Ran Apply in dir: . workspace: default

null_resource.demo: Creating...
null_resource.demo: Creation complete after 0s (ID: 4542221565395344699)

Apply complete! Resources: 1 added, 0 changed, 0 destroyed.
```

Below the comment, there's a success message with a wrench icon:

Pull request successfully merged and closed  
You're all set—the atlantis-workflow--- branch can be safely deleted.

A 'Delete branch' button is visible to the right of the message.

Atlantis was built as an alternative to Terraform Cloud automation.  
The creators of Atlantis now work at HashiCorp and is maintained by HashiCorp.

# CDK for Terraform

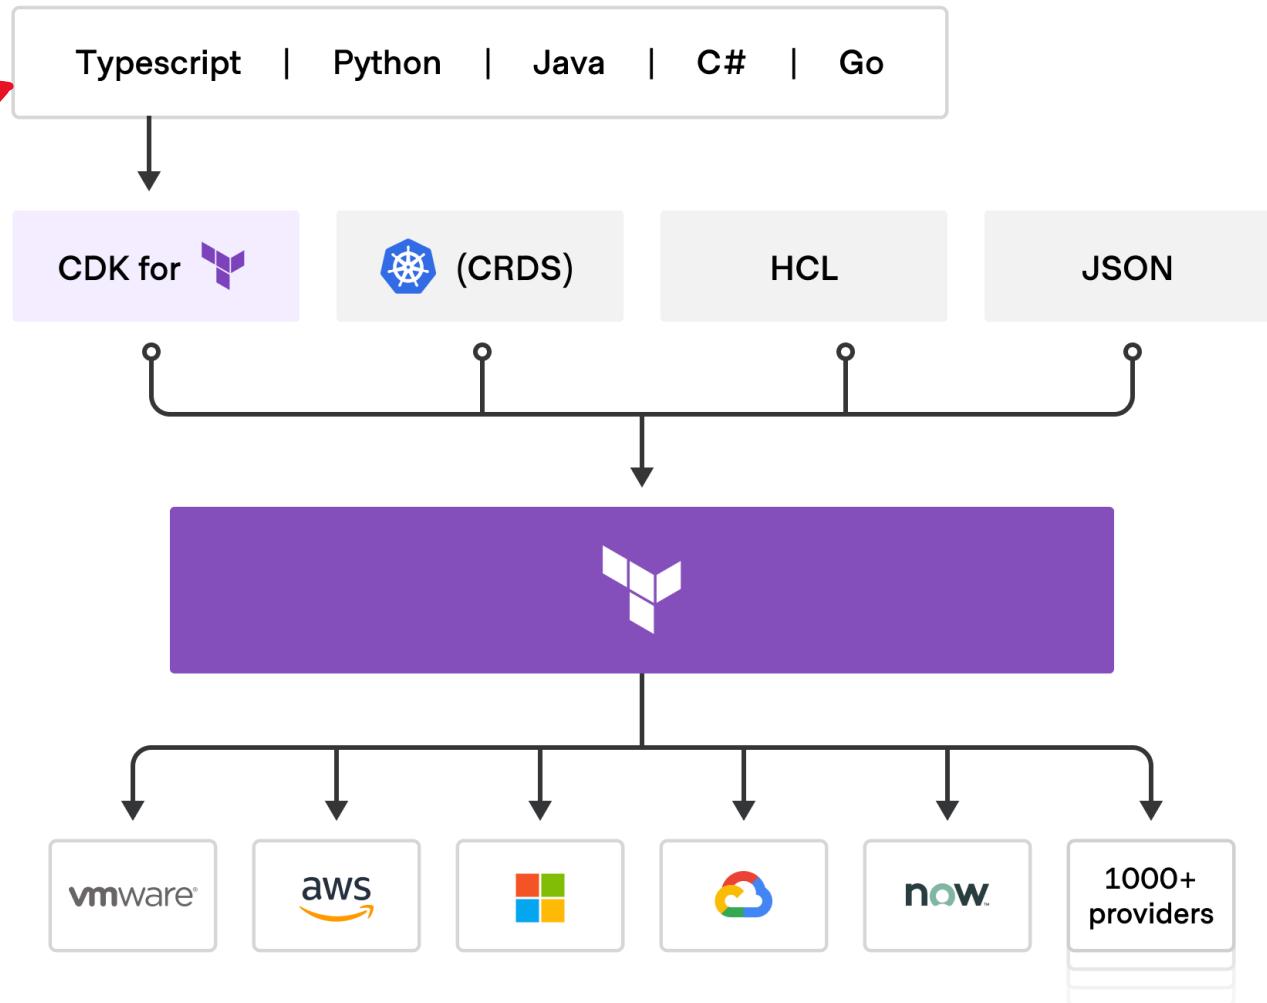


AWS Cloud Development Kit (CDK) is an **imperative** Infrastructure as Code (IaC) tool  
With SDKs for your **favorite language**

AWS CDK is intended only for AWS cloud resources.  
CDK generates out CloudFormation (CFN) templates  
(known as synthesizing) and uses that for IaC.

CDK for Terraform is a standalone project by HashiCorp that allows you to use CDK, but instead of CFN templates it generates out Terraform templates.

This allows you **to use the CDK tooling to define IaC resources for any provider** available on Terraform via CDK.



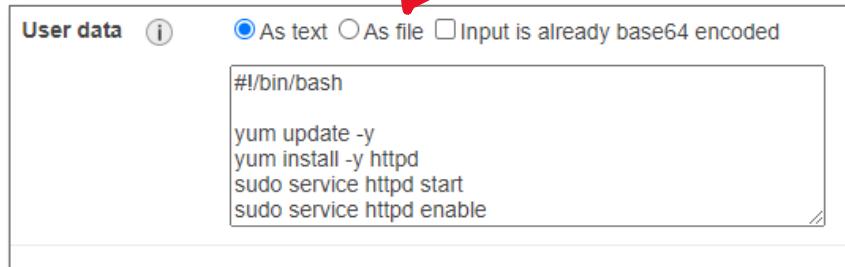
# Terraform Provisioners

**Terraform Provisioners** install software, edit files, and provision machines created with Terraform

Terraform allows you to work with two different provisioners:



Cloud-Init is an industry standard for cross-platform cloud instance initializations. When you launch a VM on a Cloud Service Provider (CSP) you'll provide a YAML or Bash script.

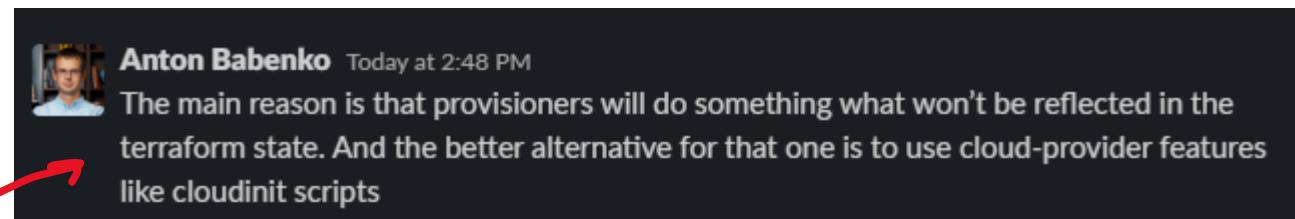


User data (i)  As text  As file  Input is already base64 encoded

```
#!/bin/bash
yum update -y
yum install -y httpd
sudo service httpd start
sudo service httpd enable
```

A screenshot of the AWS Lambda configuration interface showing the 'User data' section. It includes radio buttons for 'As text' (selected), 'As file', and 'Input is already base64 encoded'. Below is a code editor containing a Bash script to update yum, install httpd, and start the service.

Packer is an automated image-builder service. You provide a configuration file to create and provision the machine image and the image is delivered to a repository for use.



Anton Babenko Today at 2:48 PM  
The main reason is that provisioners will do something what won't be reflected in the terraform state. And the better alternative for that one is to use cloud-provider features like cloudinit scripts

A screenshot of a Slack message from Anton Babenko. The message discusses the limitations of provisioners and suggests using cloud-provider features like cloudinit scripts instead.

Provisioners should only be used as a **last resort**. For most common situations there are better alternatives.

# Terraform Provisioners

Create your own Cloud-Init script

```
users:
  - default
  - name: terraform
    gecos: terraform
    primary_group: hashicorp
    sudo: ALL=(ALL) NOPASSWD:ALL
    groups: users, admin
    ssh_import_id:
    lock_passwd: false
    ssh_authorized_keys:
      - # Paste your created SSH key here
package_upgrade: yes
package_update: yes
packages:
  - httpd
runcmd:
  - sudo service httpd start
  - sudo service httpd enable
```

Define the template file

```
data "template_file" "user_data" {
  template = file("../scripts/add-ssh-web-app.yaml")
}

resource "aws_instance" "web" {
  ami                               = data.aws_ami.ubuntu.id
  instance_type                     = "t2.micro"
  subnet_id                         = aws_subnet.subnet_public.id
  vpc_security_group_ids           = [aws_security_group.sg_22_80.id]
  associate_public_ip_address      = true
  user_data                         = data.template_file.user_data.rendered

  tags = {
    Name = "Learn-CloudInit"
  }
}
```

Reference it in the userdata for the Virtual Machine

# Terraform Provisioners

Terraform used to directly support third-party Provisioning tools in the Terraform language. Support was **deprecated** because Terraform considered using Provisioners to be poor practice suggesting better alternatives.



Cloud-Init supports Chef and Puppet, so you can just use Cloud-Init

```
puppet:
  install: true
  version: "7.7.0"
  install_type: "packages"
  collection: "puppet7"
  cleanup: true
  aio_install_url:
  "https://raw.githubusercontent.com/puppetlabs/install-
  puppet/main/install.sh"
  conf_file: "/etc/puppet/puppet.conf"
  ssl_dir: "/var/lib/puppet/ssl"
  csr_attributes_path: "/etc/puppet/csr_attributes.yaml"
  package_name: "puppet"
  exec: false
  exec_args: ['--test']
  conf:
    agent:
      server: "puppetserver.example.org"
```



ANSIBLE

It is uncertain if this advice extends to Ansible.  
Terraform and Ansible have lots of learning materials  
on how their technologies complement each other.

# Local-exec

Local-exec allows you to execute **local commands** after a resource is provisioned.

The machine that is executing Terraform *eg. terraform apply* is **where** the command will execute.

A local environment could be.....



## Local Machine

Your laptop / workstation

## Example Use Case

After you provision a VM you need to supply the Public IP to a third-party security service to add the VM IP address and you accomplish this by using locally installed third-party CLI on your build server.



## Build Server

eg. GCP Cloud Build, AWS CodeBuild, Jenkins

## Outputs vs Local-Exec

Terraform outputs allows you to output results after running Terraform apply



## Terraform Cloud Run Environment

single-use Linux virtual machine

local exec allows you to run any arbitrary commands on your local machine. Commonly used to trigger Configuration Management eg. Ansible, Chef, Puppet

# Local-exec

## command (required)

The command you want to execute

```
resource "aws_instance" "web" {  
    # ...  
  
    provisioner "local-exec" {  
        command = "echo ${self.private_ip} >> private_ips.txt"  
    }  
}
```

## working\_dir

Where the command will be executed  
eg. /user/andrew/home/project

```
resource "null_resource" "example2" {  
    provisioner "local-exec" {  
        command = "Get-Date > completed.txt"  
        interpreter = ["PowerShell", "-Command"]  
    }  
}
```

## interpreter

The entry point for the command.  
What local program will run the command eg.  
Bash, Ruby, AWS CLI, PowerShell

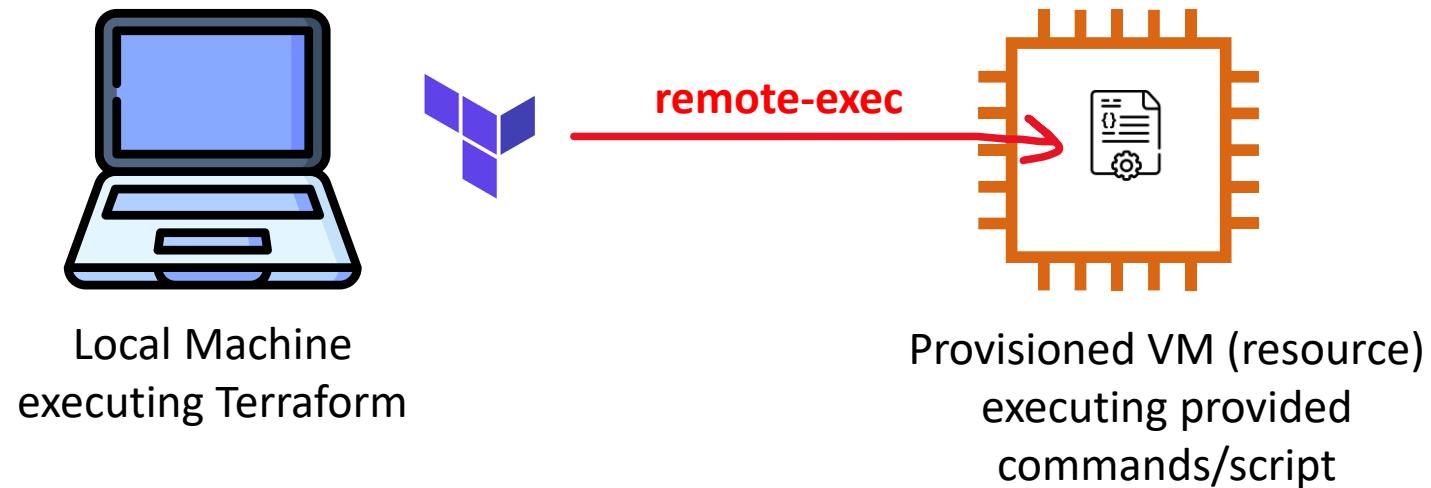
## Environment

Key and value pair of environment variables

```
resource "aws_instance" "web" {  
    # ...  
  
    provisioner "local-exec" {  
        command = "echo $KEY $SECRET >> credentials.yml"  
  
        environment = {  
            KEY = "Jlmyn61Qmc8wiux9zprS9v4n"  
            SECRET = "U32KIy2T1fW3xGD1FLggXjuN"  
        }  
    }  
}
```

# Remote-exec

Remote-exec allows you to execute **commands on a target resource** *after* a resource is provisioned.



Remote-Exec is useful for provisioning a Virtual Machine with a simple set of commands.

For more complex tasks its recommended to use Cloud-Init, and strongly recommended in all cases to bake Golden Images via Packer or EC2 Image Builder

# Remote-exec

**Remote Command has three different modes:**

**Inline** - list of command strings

**Script** - relative or absolute local script that will be copied to the remote resource and then executed

**Scripts** - relative or absolute local scripts that will be copied to the remote resource and then executed and executed in order.

*You can only choose to use one mode at a time*

```
resource "aws_instance" "web" {
    # ...

    provisioner "remote-exec" {
        inline = [
            "puppet apply",
            "consul join ${aws_instance.web.private_ip}",
        ]
    }
}
```

```
resource "aws_instance" "web" {
    # ...

    provisioner "remote-exec" {
        scripts = [
            "./setup-users.sh",
            "/home/andrew/Desktop/bootstrap"
        ]
    }
}
```

# File

**file provisioner** is used to copy files or directories from our local machine to the newly created resource

**Source** – the local file we want to upload to the remote machine

**Content** – a file or a folder

**Destination** – where you want to upload the file on the remote machine

You may require a connection block within the provisioner for authentication

```
resource "aws_instance" "web" {
    # ...

    # Copies the myapp.conf file to /etc/myapp.conf
    provisioner "file" {
        source      = "conf/myapp.conf"
        destination = "/etc/myapp.conf"
    }

    # Copies the string in content into /tmp/file.log
    provisioner "file" {
        content      = "ami used: ${self.ami}"
        destination = "/tmp/file.log"
    }

    # Copies the configs.d folder to /etc/configs.d
    provisioner "file" {
        source      = "conf/configs.d"
        destination = "/etc"
    }

    # Copies all files and folders in apps/app1 to D:/IIS/webapp1
    provisioner "file" {
        source      = "apps/app1/"
        destination = "D:/IIS/webapp1"
    }
}
```

# Connection

A connection block tells a **provisioner** or **resource** how to establish a connection

You can connect via **SSH**

With SSH you can connect through a Bastion Host eg:

- bastion\_host
- bastion\_host\_key
- bastion\_port
- bastion\_user
- bastion\_password
- bastion\_private\_key
- bastion\_certificate

You can connect via **Windows Remote Management (winrm)**

```
# Copies the file as the root user using SSH
provisioner "file" {
  source      = "conf/myapp.conf"
  destination = "/etc/myapp.conf"

  connection {
    type      = "ssh"
    user      = "root"
    password = "${var.root_password}"
    host      = "${var.host}"
  }
}

# Copies the file as the Administrator user using WinRM
provisioner "file" {
  source      = "conf/myapp.conf"
  destination = "C:/App/myapp.conf"

  connection {
    type      = "winrm"
    user      = "Administrator"
    password = "${var.admin_password}"
    host      = "${var.host}"
  }
}
```

# Null Resources

**null\_resource** is a placeholder for resources that have no specific association to a provider resources.

You can provide a **connection** an  
**triggers** to a resource

**Triggers** is a map of values which **should**  
**cause this set of provisioners to re-run.**

Values are meant to be interpolated  
references to variables or attributes of  
other resources

```
resource "aws_instance" "cluster" {
  count = 3

  # ...
}

resource "null_resource" "cluster" {
  # Changes to any instance of the cluster requires re-provisioning
  triggers = {
    cluster_instance_ids = "${join(", ", aws_instance.cluster.*.id)}"
  }

  # Bootstrap script can run on any instance of the cluster
  # So we just choose the first in this case
  connection {
    host = "${element(aws_instance.cluster.*.public_ip, 0)}"
  }

  provisioner "remote-exec" {
    # Bootstrap script called with private_ip of each node in the cluster
    inline = [
      "bootstrap-cluster.sh ${join(" ", aws_instance.cluster.*.private_ip)}",
    ]
  }
}
```

# Terraform Data

Similar to null\_resources but **does not require or the configuration of a provider**.

```
resource "null_resource" "main" {  
    triggers = {  
        version = var.version  
    }  
  
    provisioners "local-exec" {  
        command = "echo ${self.triggers.version}"  
    }  
}
```

```
resource "terraform_data" "main" {  
    triggers_replace = [  
        version  
    ]  
  
    provisioners "local-exec" {  
        command = "echo ${self.triggers_replace}"  
    }  
}
```



# Terraform Data – replace\_triggered\_by

```
variable "revision" {
  default = 1
}

resource "terraform_data" "replacement" {
  input = var.revision
}

resource "example_database" "prod" {
  lifecycle {
    replace_triggered_by = [terraform_data.replacement]
  }
}
```

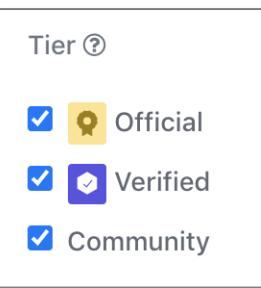
# Terraform Providers



Providers are Terraform Plugins that allow you to interact with:

- Cloud Service Providers (CSPs) eg. **AWS, Azure, GCP**
- Software as a Service (SaaS) Providers eg. **Github, Angolia, Stripe**
- Other APIs eg. **Kubernetes, Postgres**

Providers are required for your Terraform Configuration file to work.



Providers come in three tiers:

Official — Published by the company that owns the provider technology or service

Verified — actively maintained, up-to-date and compatible with both Terraform and Provider

Community — published by a community member but no guarantee of maintenance, up-to-date or compatibility

Providers are distributed separately from Terraform and the plugins must be downloaded before use.

**terraform init** will download the necessary provider plugins listed in a Terraform configuration file.

# Terraform Registry

**Terraform Registry** is a **website portal** to **browse, download or publish** available **Providers** or **Modules**  
<https://registry.terraform.io>



A screenshot of the Terraform Registry website. The top navigation bar includes the HashiCorp logo, a search bar labeled "Search Providers and Modules", and links for "Browse", "Publish", and "Sign-in". Below the navigation is a main content area with tabs for "Providers" (which is selected) and "Modules". On the left, there's a sidebar with "FILTERS" (set to "Official", "Verified", and "Community") and a "Category" list. The main content area shows a grid of provider icons: AWS (orange), Azure (blue), Google Cloud Platform (blue), Kubernetes (blue), Oracle Cloud Infrastructure (red), and Alibaba Cloud (dark grey). At the bottom, there are three small preview cards for Active Directory, Archive, and Azure Active Directory.

## Provider

A provider is a plugin that is mapping to a Cloud Service Provider (CSPs) API.

## Module

A module is a group of configuration files that provide common configuration functionality.

- Enforces best practices
- reduce the amount of code
- Reduce time to develop scripts

Everything published to Terraform Registry is **public-facing**

# Terraform Registry – Providers

You can easily find **documentation** and **code sample** for providers

The screenshot shows the Terraform Registry provider page for the AWS provider. At the top, there's a navigation bar with 'Providers / hashicorp / aws / Version 3.53.0' and a 'Latest Version' button. Below the navigation, there's a header with the provider name 'aws' and its official status. A 'Public Cloud' badge is also present. The main content area contains a brief description of the provider's functionality, its version (3.53.0), and download statistics (2 days ago, 430.4M). A large callout box highlights the 'Documentation' tab and the 'USE PROVIDER' button. Inside the callout box, there's a section titled 'How to use this provider' with Terraform code examples.

```
terraform {  
    required_providers {  
        aws = {  
            source = "hashicorp/aws"  
            version = "3.53.0"  
        }  
    }  
  
    provider "aws" {  
        # Configuration options  
    }  
}
```

The screenshot shows the documentation page for the AWS App Mesh gateway route provider. At the top, there's a navigation bar with 'Overview' and 'Documentation'. The main content area starts with a 'Resource:' section for 'aws\_appmesh\_gateway\_route'. It provides a brief description and example usage. Below this, there's a sidebar with a tree view of AWS services, where 'AppMesh' is selected. The 'Resources' section lists various resources, with 'aws\_appmesh\_gateway\_route' being highlighted.

**Resource:**  
**aws\_appmesh\_gateway\_route**

Provides an AWS App Mesh gateway route resource.

**Example Usage**

```
resource "aws_appmesh_gateway_route" "example" {  
    name          = "example-gateway-route"  
    mesh_name     = "example-service-mesh"  
    virtual_gateway_name = aws_appmesh_virtual_gateway.example  
  
    spec {  
        http_route {  
            action {  
                target {  
                    virtual_service {  
                        virtual_service_name = aws_appmesh_virtual_node.example  
                    }  
                }  
            }  
        }  
    }  
}
```

# Terraform Registry

A list of dependent modules

The screenshot shows a Terraform module page for the AWS VPC module. At the top left is the AWS logo. Next to it is the module name "vpc" followed by a blue hexagon icon with a white checkmark. Below the name is the text "AWS". A description follows: "Terraform module which creates VPC resources on AWS". To the right is a dropdown menu showing "Version 3.2.0 (latest)". On the left side of the main content area, there's a red curved arrow pointing from the "A list of dependent modules" text to the "Submodules" button. At the bottom of the main content area, there are two buttons: "Submodules" and "Examples". A red arrow points from the "Lots of examples for common use cases" text to the "Examples" button. On the right side, there's a callout box with the heading "Provision Instructions" containing sample Terraform code. A red arrow points from the "You can easily grab the module code:" text to the "Provision Instructions" box.

Lots of examples for common use cases

You can easily grab the **module code**:

AWS

vpc

AWS

Terraform module which creates VPC resources on AWS

Published June 28, 2021 by [terraform-aws-modules](#)

Module managed by [antonbabenko](#)

Total provisions: 10.5M

Source Code: [github.com/terraform-aws-modules/terraform-aws-vpc](https://github.com/terraform-aws-modules/terraform-aws-vpc) (report an issue)

Submodules Examples

Version 3.2.0 (latest) ▾

**Provision Instructions**

Copy and paste into your Terraform configuration, insert the variables, and run `terraform init` :

```
module "vpc" {  
  source = "terraform-aws-modules/vpc/  
  version = "3.2.0"  
  # insert the 19 required variables !  
}
```

Readme Inputs (160) Outputs (108) Dependency (1) Resources (76)

# Terraform Cloud – Private Registry

Terraform Cloud allows you to **publish private modules** for your Organization within the **Terraform Cloud Private Registry**

When creating a module you need to connect to a Version Control System (VCS) and choose a repository

## Add Module

This module will be created under the current organization, **ExamPro**. Modules can be added from all supported VCS providers. [🔗](#)

1 Connect to VCS

2 Choose a repository

3 Confirm selection

### Connect to a version control provider

Choose the version control provider that hosts your module source code.



# Terraform Providers Command

Get a list of the current providers you are using

```
terraform providers
Providers required by configuration:
.
├── provider[registry.terraform.io/hashicorp/azurerm] 2.72.0
├── provider[registry.terraform.io/hashicorp/google] 3.80.0
└── provider[registry.terraform.io/hashicorp/aws] 3.54.0
└── module.network
    └── provider[registry.terraform.io/hashicorp/azurerm]
└── module.linuxservers
    ├── provider[registry.terraform.io/hashicorp/azurerm]
    ├── provider[registry.terraform.io/hashicorp/random]
    └── module.os
```

# Terraform Provider Configuration

How to reference  
an alias provider



```
provider "aws" {  
  alias = "west"  
  region = "us-west-2"  
}
```

Set an alternative provider



```
resource "aws_instance" "foo" {  
  provider = aws.west  
  # ...  
}
```

How to set alias provider  
for a parent module



```
terraform {  
  required_providers {  
    mycloud = {  
      source  = "mycorp/mycloud"  
      version = "~> 1.0"  
      configuration_aliases = [ mycloud.alternate ]  
    }  
  }  
}
```

How to set alias provider  
for a child module



```
module "aws_vpc" {  
  source = "./aws_vpc"  
  providers = {  
    aws = aws.west  
  }  
}
```

# Terraform Modules

A Terraform module is a group of configuration files that provide common configuration functionality.

- Enforces best practices
- reduce the amount of code
- Reduce time to develop scripts

## AWS Provider (not a module)

If you had to create a VPC you would have specific many networking resources.

```
resource "aws_vpc" "main" {  
  cidr_block      = "10.0.0.0/16"  
  instance_tenancy = "default"  
  
  tags = {  
    Name = "main"  
  }  
  
resource "aws_subnet" "main" {  
  vpc_id      = aws_vpc.main.id  
  cidr_block  = "10.0.1.0/24"  
  
  tags = {  
    Name = "Main"  
  }  
  
  # ...
```

## AWS VPC Module

Using a module you can use a a shorthand Domain Specific Language (DSL)  
That will reduce the amount of work.

```
module "vpc" {  
  source = "terraform-aws-modules/vpc/aws"  
  
  name      = "my-vpc"  
  cidr     = "10.0.0.0/16"  
  
  azs        = ["eu-west-1a", "eu-west-1b", "eu-west-1c"]  
  private_subnets = ["10.0.1.0/24", "10.0.2.0/24", "10.0.3.0/24"]  
  public_subnets  = ["10.0.101.0/24", "10.0.102.0/24", "10.0.103.0/24"]  
  
  enable_nat_gateway = true  
  enable_vpn_gateway = true  
  
  tags = {  
    Terraform = "true"  
    Environment = "dev"  
  }  
}
```



Modules: Imagine clicking a wizard that creates many cloud resources e.g. VPC Wizard

# Terraform Modules

# Azure VM via the **Azure Provider**

```

resource "azurerm_resource_group" "hytterformancegroup" {
  name     = "hytterformancegroup"
  location = "EastUS"
}

tags = [
  environment = "Hytterformance"
]

r_create_virtual_network "hytterformance_vnet" {
  name                = "hytterformancevnet"
  address_space       = ["10.0.0.0/16"]
  location            = "EastUS"
  resource_group_name = azurerm_resource_group.hytterformancegroup.name
  subnet_prefixes     = ["10.0.1.0/24", "10.0.2.0/24"]
  tags                = [
    environment = "Hytterformance"
  ]
}

resource "azurerm_subnet" "hytterformancesubnet" {
  name                 = "hytterformancesubnet"
  resource_group_name = azurerm_resource_group.hytterformancegroup.name
  virtual_network_name = azurerm_virtual_network.hytterformance_vnet.name
  address_prefixes     = ["10.0.1.0/24", "10.0.2.0/24"]
  tags                = [
    environment = "Hytterformance"
  ]
}

resource "azurerm_public_ip" "hytterformancepublicip" {
  name                = "hytterformancepublicip"
  allocation_method   = "Dynamic"
  location            = "EastUS"
  resource_group_name = azurerm_resource_group.hytterformancegroup.name
  tags                = [
    environment = "Hytterformance"
  ]
}

resource "azurerm_network_security_group" "hytterformance" {
  name                = "hytterformance"
  location            = "EastUS"
  resource_group_name = azurerm_resource_group.hytterformancegroup.name
  tags                = [
    environment = "Hytterformance"
  ]
}

resource "azurerm_nsg_rule" "hytterformanceinboundrule" {
  name                = "hytterformanceinboundrule"
  direction           = "Inbound"
  priority            = 1000
  protocol            = "Tcp"
  source_port_range   = "22"
  destination_port_range = "22"
  source_address_prefix = "*"
  destination_address_prefix = "*"
  destination_port_range = "*"
  tags                = [
    environment = "Hytterformance"
  ]
}

resource "azurerm_nsg_rule" "hytterformanceoutboundrule" {
  name                = "hytterformanceoutboundrule"
  direction           = "Outbound"
  priority            = 1000
  protocol            = "Tcp"
  source_port_range   = "22"
  destination_port_range = "22"
  source_address_prefix = "*"
  destination_address_prefix = "*"
  tags                = [
    environment = "Hytterformance"
  ]
}

r_create_storage_interface "hytterformance" {
  resource_group_name = azurerm_resource_group.hytterformancegroup.name
  network_interface_id = azurerm_network_interface.hytterformance.id
  storage_account_id = azurerm_storage_account.hytterformance.id
}

resource "azurerm_storage_account" "hytterformance" {
  name                = "hytterformance"
  location            = "EastUS"
  account_tier        = "Standard"
  account_replication_type = "LRS"
  tags                = [
    environment = "Hytterformance"
  ]
}

r_update_the_identity_for_the_storage_interface "hytterformance" {
  resource_group_name = azurerm_resource_group.hytterformancegroup.name
  storage_account_id = azurerm_storage_account.hytterformance.id
  network_interface_id = azurerm_network_interface.hytterformance.id
}

resource "azurerm_key_vault" "hytterformancekeyvault" {
  name                = "hytterformancekeyvault"
  location            = "EastUS"
  resource_group_name = azurerm_resource_group.hytterformancegroup.name
  tags                = [
    environment = "Hytterformance"
  ]
}

resource "azurerm_key_vault_key" "hytterformancekey" {
  name                = "hytterformancekey"
  key_vault_name       = azurerm_key_vault.hytterformancekeyvault.name
  key_type            = "RSA"
  key_size            = 2048
  key_usage           = "EncryptAndDecrypt"
  enable_purge_protection = true
  tags                = [
    environment = "Hytterformance"
  ]
}

resource "azurerm_virtual_machine" "hytterformancevm" {
  name                = "hytterformancevm"
  location            = "EastUS"
  resource_group_name = azurerm_resource_group.hytterformancegroup.name
  network_interface_id = azurerm_network_interface.hytterformance.id
  size                = "Standard_B1ms"
  os_disk {
    name                = "hytterformanceosdisk"
    storage_account_url = azurerm_storage_account.hytterformance.id
    storage_account_type = "Premium_LRS"
  }
  source_image_reference {
    publisher          = "Canonical"
    offer               = "UbuntuServer"
    sku                = "16.04-LTS"
    version            = "latest"
  }
  computer_name       = "hytterformancevm"
  disable_password_authentication = true
  admin_username      = "hytterformance"
  admin_password      = file("hytterformancevm/admin_password.txt").base64
}

host_discovery {
  storage_account_url = azurerm_storage_account.hytterformancegroup.primary_blob_endpoint
}

tags = [
  environment = "Hytterformance"
]

```

## Azure VM via a **Compute and Network Module**

```
resource "azurerm_resource_group" "example" {
  name      = "example-resources"
  location  = "West Europe"
}

module "linuxservers" {
  source          = "Azure/compute/azurerm"
  resource_group_name = azurerm_resource_group.example.name
  vm_os_simple    = "UbuntuServer"
  public_ip_dns    = ["linsimplevmips"]
  vnet_subnet_id   = module.network.vnet_subnets[0]

  depends_on = [azurerm_resource_group.example]
}

module "network" {
  source          = "Azure/network/azurerm"
  resource_group_name = azurerm_resource_group.example.name
  subnet_prefixes  = ["10.0.1.0/24"]
  subnet_names     = ["subnet1"]

  depends_on = [azurerm_resource_group.example]
}

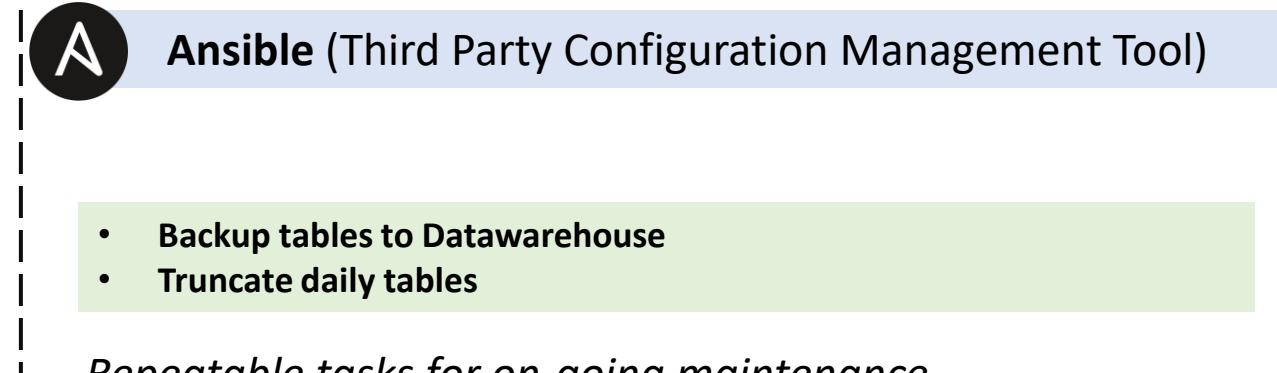
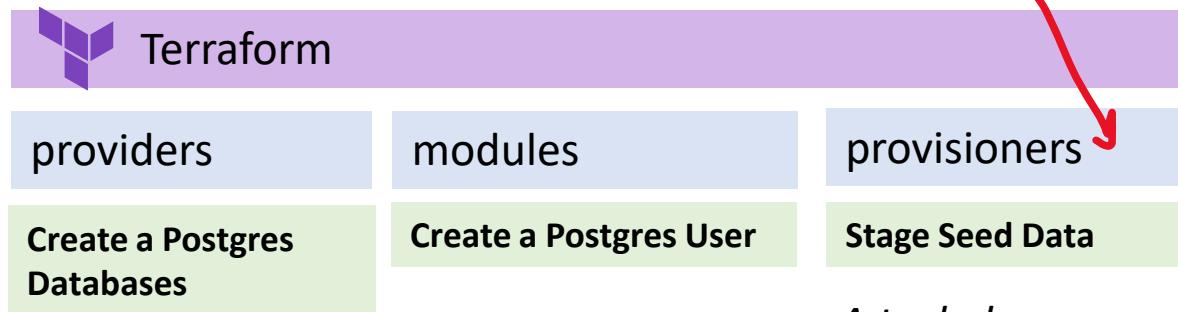
output "linux_vm_public_name" {
  value = module.linuxservers.public_ip_dns_name
}
```

# Terraform Providers – The Fine Line

**The Fine Line** is understanding **the granularities of responsibility** between Terraform Infrastructure as Code and Third-Party Configuration Management.

When you have a Postgres Database which is a typical resource compared to a cloud service like an Amazon S3  
Who should automate what?

A Terraform Provisioner could  
be using Ansible



*Entities. If you want governance or asset resource management*

*A task done one time to setup the database*

*Repeatable tasks for on-going maintenance*

# Terraform Language

Terraform files contain the configuration information about **providers** and **resources**.

Terraform files end in the extension of **.tf** or either **.tf.json**

Terraform files are written in the **Terraform Language** and is the extension of HCL

Terraform language consists of only a few basic elements:

- **Blocks** — containers for other content, represent an object
  - block type — can have zero or more labels and a body
  - block label — name of a block
- **Arguments** — assign a value to a name
  - They appear within blocks
- **Expressions** — represent a value, either literally or by referencing and combining other values
  - They appear as values for arguments, or within other expressions.

```
resource "aws_vpc" "main" {  
    cidr_block = var.base_cidr_block  
}  
  
<BLOCK TYPE> "<BLOCK LABEL>" "<BLOCK LABEL>" {  
    # Block body  
    <IDENTIFIER> = <EXPRESSION> # Argument  
}
```

You might come across HashiCorp Configuration Language (HCL) this is the low-level language for both the Terraform Language and alternative JSON syntax

# Hashicorp Configuration Files – Syntax

## Arguments

An *argument* assigns a value to a particular name

- arguments appear within blocks

```
image_id = "abc123"
```

## Blocks

A block is a container for other content

Block has a type “**resource**”  
and it expects two Block Labels

different types have different number of  
required labels.

```
resource "aws_instance" "example"  
{  
    ami = "abc123"  
  
    network_interface {  
        # ...  
    }  
}
```

A block can contain  
**nested** blocks

The brackets { } indicate the block **body**

# Hashicorp Configuration Files – Syntax

```
image_id = "abc123"
```



**Identifiers**, uniquely identify objects of constructs in a Terraform Language:

- Argument names
- Block type names
- names Terraform-specific constructs eg. resources, input variables

Identifiers can contain:

- letters eg. abc
- digits eg. 012
- underscores eg. \_
- hyphens eg. -

The first letter of an Identifiers cannot be a digit

# Hashicorp Configuration Files – Syntax

Terraform language supports **3 different syntaxes** for comments

```
# this is also a single line comment
```

```
// this is a single line comment
```

```
/*
This
is a
multiple comment
*/
```



This is the default comment style.  
Formatting tools may transform // to #  
eg. terraform fmt

Terraform configuration files must always be **UTF-8 encoded**

Terraform accepts configuration files with either Unix-style line endings (LF) or Windows-style line endings (CR + LF)

# Hashicorp Configuration Files – Alternate JSON Syntax

Terraform also supports an *alternative syntax* that is JSON-compatible  
Terraform expects JSON syntax files to be named with **.tf.json**

This syntax is useful when generating portions  
of a configuration programmatically, since  
existing JSON libraries can be used to prepare  
the generated configuration files.

Example of **JSON syntax**

```
{  
  "resource": {  
    "aws_instance": {  
      "example": {  
        "instance_type": "t2.micro",  
        "ami": "ami-abc123"  
      }  
    }  
  }  
}
```

# Terraform Settings

The special **terraform configuration block type** eg. **terraform { ... }** is used to configure some behaviors of Terraform itself

In Terraform settings we can specify:

- **required\_version**
  - The expected version of terraform
- **required\_providers**
  - The providers that will be pull during an terraform init
- **experiments**
  - Experimental language features, that the community can try and provide feedback
- **provider\_meta**
  - module-specific information for providers

```
terraform {
  required_providers {
    aws = {
      version = ">= 2.7.0"
      source  = "hashicorp/aws"
    }
  }
}
```

# HashiCorp Configuration Language

HCL is an open-source toolkit for creating **structured configuration languages** that are both human and machine friendly, for use with command-line tools

[github.com/hashicorp/hcl](https://github.com/hashicorp/hcl)

## HashiCorp Configuration Language (HCL)

|                                                                                     |                                                      |                                                 |
|-------------------------------------------------------------------------------------|------------------------------------------------------|-------------------------------------------------|
|    | Terraform Language (.tf)                             | eg. Dynamic Blocks, For Each...                 |
|    | Packer Template (.pkr.hcl)                           |                                                 |
|    | Vault Policies (no extension)                        |                                                 |
|    | Boundary Controllers and Workers (.hcl)              |                                                 |
|    | Consul Configuration (.hcl)                          |                                                 |
|   | Waypoint Application Configuration(.hcl)             |                                                 |
|  | Nomad Job Specifications (.nomad)                    |                                                 |
|  | Shipyard Blueprint (.hcl)                            |                                                 |
|  | Sentinel Policies <span style="color: red;">X</span> | Doesn't use HCL but its own ACL custom language |

# Input Variables

**Input variables** (aka variables or Terraform Variables)  
are **parameters** for Terraform modules

You can declare variables in either:

- The root module
- The child modules

**Default** A default value which then makes the variable optional

**type** This argument specifies what value types are accepted for the variable

**Description** This specifies the input variable's documentation

**Validation** A block to define validation rules, usually in addition to type constraints

**Sensitive** Limits Terraform UI output when the variable is used in configuration

Variables are defined  
via **variable blocks**.

```
variable "image_id" {  
  type = string  
}  
  
variable "availability_zone_names" {  
  type    = list(string)  
  default = ["us-west-1a"]  
}  
  
variable "docker_ports" {  
  type = list(object({  
    internal = number  
    external = number  
    protocol = string  
  }))  
  default = [  
    {  
      internal = 8300  
      external = 8300  
      protocol = "tcp"  
    }  
  ]  
}
```

# Variable Definitions Files

A variable definitions file allows you to set the values for multiple variables at once.

Variable definition files are named .tfvars or tfvars.json

By default **terraform.tfvars** will be autoloaded when included in the root of your project directory

```
image_id = "ami-abc123"  
availability_zone_names = [  
    "us-east-1a",  
    "us-west-1c",  
]
```

Variable Definition Files use the Terraform Language.



# Variables via Environment Variables

A variable value can be defined by Environment Variables

Variable starting with **TF\_VAR\_<name>** will be read and loaded



```
export TF_VAR_image_id=ami-abc123
```

# Loading Input Variables

**Default Autoloaded Variables file**  
**terraform.tfvars**

When you create a named **terraform.tfvars** file it will be automatically loaded when running terraform apply

**Additional Variables Files (not autoloaded)**  
my\_variables.tfvars

You can create additional variables files eg. dev.tfvars, prod.tfvars  
They will not be autoloaded (you'll need to specify them in via command line)

**Additional Variables Files (autoloaded)**  
my\_variables.**auto**.tfvars

If you name your file with auto.tfvars it will always be loaded

**Specify a Variables file via Command Line**  
**-var-file** dev.tfvars

You can specify variables inline via the command line for individual overrides

**Inline Variables via Command Line**  
**-var** ec2\_type="t2.medium"

You can specify variables inline via the command line for individual overrides

**Environment Variables**  
**TF\_VAR\_my\_variable\_name**

Terraform will watch for environment variables that begin with TF\_VAR\_ and apply those as variables

# Variable Definition Precedence

You can override variables via many files and commands

The definition precedence is the order in which Terraform will read variables and as it goes down the list it will override variable.

- 
- Environment Variables
  - `terraform.tfvars`
  - `terraform.tfvars.json`
  - `*.auto.tfvars` or `*.auto.tfvars.json`
  - `-var` and `-var-file`

# Output Values

Output Values are computed values after a Terraform apply is performed.

Outputs allows you:

- to obtain information after resource provisioning e.g. public IP address
- output a file of values for programmatic integration
- Cross-reference stacks via outputs in a state file via `terraform_remote_state`

You can optionally provide a description

You can mark the output as sensitive so it does not show in output of your Terminal

```
output "db_password" {  
    value      = aws_db_instance.db.password  
    description = "The password for logging in to the database."  
    sensitive  = true  
}
```



Sensitive outputs will still be visible within the statefile.

# Output Values

To print all the outputs for a statefile  
use the **terraform output**



```
$ terraform output
lb_url = "http://lb-5YI-project-alpha-dev-2144336064.us-east-1.elb.amazonaws.com/"
vpc_id = "vpc-004c2d1ba7394b3d6"
web_server_count = 4
```

Print a specific output with  
**terraform output <name>**



```
$ terraform output lb_url
"http://lb-5YI-project-alpha-dev-2144336064.us-east-1.elb.amazonaws.com/"
```

Use the **-json** flag to get  
output as json data.



```
$ terraform output -json
{
  "db_password": {
    "sensitive": true,
    "type": "string",
    "value": "notasecurepassword"
  },
}
```

Use the **-raw** flag to  
preserve quotes for strings



```
$ curl $(terraform output -raw lb_url)
<html><body><div>Hello, world!</div></body></html>
```

# Local Values

A local value (locals) **assigns a name to an expression**, so you can **use it multiple times within a module** without repeating it.

Locals are set using  
the **locals block**

You can define  
**multiple** locals blocks

You can **reference**  
**locals within locals**

```
locals {  
    service_name = "forum"  
    owner        = "Community Team"  
}  
  
locals {  
    # IDs for multiple sets of EC2 instances, merged together  
    instance_ids = concat(aws_instance.blue.*.id, aws_instance.green.*.id)  
}  
  
locals {  
    # Common tags to be assigned to all resources  
    common_tags = {  
        Service = local.service_name  
        Owner   = local.owner  
    }  
}
```

Static value  
computed values

# Local Values

Once a local value is declared, you can reference it in expressions as `local.<NAME>`.

```
resource "aws_instance" "example" {  
    # ...  
  
    tags = local.common_tags  
}
```



When you referencing **you use the singular “local”**

Locals help can help DRY up your code.

It is best practice to **use locals sparingly** since it Terraform is intended to be declarative and overuse of locals can make it difficult to determine what the code is doing.

# Data Sources

**Data sources** allow Terraform **use information defined outside of Terraform**, defined by another separate Terraform configuration, or modified by functions.

You specific what kind of external resource you want to select

```
data "aws_ami" "web" {  
    filter {  
        name  = "state"  
        values = ["available"]  
    }  
  
    filter {  
        name  = "tag:Component"  
        values = ["web"]  
    }  
  
    most_recent = true  
}
```

You use filters to narrow down the selection

```
resource "aws_instance" "web" {  
    ami           = data.aws_ami.web.id  
    instance_type = "t1.micro"  
}
```

You use **data.** to reference data sources

# References to Named Values

Named Values are **built-in expressions** to **reference various values** such as:

Resources <Resource Type>.<Name> e.g. aws\_instance.my\_server

Input variables **var.<Name>**

Local values **local.<Name>**

Child module outputs **module.<Name>**

Data sources **data.<Data Type>.<Name>**

Filesystem and workspace info

- path.module - path of the module where the expression is placed
- path.root - path of the root module of the configuration
- path.cwd - path of the current working directory
- terraform.workspace – name of the currently selected workspace

Block-local values (within the body of blocks)

- count.**index** (when you use the count meta argument)
- each.**key** / each.**value** (when you use the for\_each meta argument )
- **self.<attribute>** - self reference information within the block (provisioners and connections)

Named values resemble the attribute notation for map (object) values but are not objects and do not act as objects. You cannot use square brackets to access attribute of Named Values like an object.

# Resource Meta Arguments

Terraform language defines several meta-arguments, which can be used with any **resource type** to change the behavior of resources.

- **depends\_on**, for specifying explicit dependencies
- **count**, for creating multiple resource instances according to a count
- **for\_each**, to create multiple instances according to a map, or set of strings
- **provider**, for selecting a non-default provider configuration
- **lifecycle**, for lifecycle customizations
- **provisioner** and connection, for taking extra actions after resource creation

# *depends\_on*

The order of which resources are provisioned is important when resources depend on others before they are provisioned.

Terraform implicitly can determine the order of provision for resources but there may be some cases where it cannot determine the correct order.

```
resource "aws_iam_role_policy" "example" {
    name      = "example"
    role      = aws_iam_role.example.name
    # ...
}

resource "aws_instance" "example" {
    ami          = "ami-a1b2c3d4"
    instance_type = "t2.micro"
    iam_instance_profile = aws_iam_instance_profile.example
    depends_on = [
        aws_iam_role_policy.example,
    ]
}
```

**depends\_on** allows you to explicitly specify a dependency of a resource.



# count

When you are managing a pool of objects eg. a fleet of Virtual Machines you can use **count**.

Specify the **amount** of instances you want

```
resource "aws_instance" "server" {  
    count = 4 # create four similar EC2 instances  
  
    ami           = "ami-a1b2c3d4"  
    instance_type = "t2.micro"  
  
    tags = {  
        Name = "Server ${count.index}"  
    }  
}
```

Get the current count **value**  
(index) via **count.index**

This value starts at **0**

Count can accept **numeric expressions**:

- Must be whole number
- Number must be known before configuration

```
resource "aws_instance" "server" {  
    # Create one instance for each subnet  
    count = length(var.subnet_ids)  
    # ....
```

# for\_each

**for\_each** is similar to count for managing multiple related objects.  
But you can iterate over a map for more dynamic values.

```
resource "azurerm_resource_group" "rg" {  
  for_each = {  
    a_group = "eastus"  
    another_group = "westus2"  
  }  
  name      = each.key  
  location = each.value  
}
```

With a map:

**each.key** – print out the current key

**each.value** – print out the current value

```
resource "aws_iam_user" "the-accounts" {  
  for_each = toset( ["Todd", "James", "Alice", "Dottie"] )  
  name      = each.key  
}
```

With a list:

**each.key** – print out the current key

# Resource Behaviour

When you execute an execution order via Terraform Apply it will perform one of the following to a resource:

## Create

- resources that exist in the configuration but are not associated with a real infrastructure object in the state.

## Destroy

- resources that exist in the state but no longer exist in the configuration.

## Update in-place

- resources whose arguments have changed.

## Destroy and re-create

- resources whose arguments have changed but which cannot be updated in-place due to remote API limitations.

```
Terraform used the selected provider to create
```

```
Terraform will perform the following actions:
```

```
# aws_instance.my_example_server
+ resource "aws_instance" "my_example_server" {
    + ami
```

```
Terraform used the selected provider to destroy
```

```
Terraform will perform the following actions:
```

```
# aws_instance.my_example_server
- resource "aws_instance" "my_example_server" {
    - ami
```

```
Terraform used the selected providers to generate the configuration changes:
~ update in-place
```

```
Terraform will perform the following actions:
```

```
# aws_instance.my_example_server will be replaced
-/+ resource "aws_instance" "my_example_server" {
    ~ arn
```

```
Terraform used the selected providers to generate the configuration changes:
-/+ destroy and then create replacement
```

```
Terraform will perform the following actions:
```

```
# aws_instance.my_example_server will be replaced
-/+ resource "aws_instance" "my_example_server" {
    ~ arn
```

# lifecycle

**Lifecycle block** allows you to change what happens to resource e.g. create, update, destroy.  
Lifecycle blocks are nested within resources

## **create\_before\_destroy** (bool)

When replacing a resource first create the new resource before deleting it (the default is destroy old first)

## **prevent\_destroy** (bool)

Ensures a resource is not destroyed

## **ignore\_changes** (list of attributes)

Don't change the resource (create, update, destroy) the resource if a change occurs for the listed attributes.

```
resource "azurerm_resource_group" "example" {
    # ...

    lifecycle {
        create_before_destroy = true
    }
}
```

# Resource Providers and Alias

If you need to override the default provider for a resource you can create alternative provider with **alias**

```
# default configuration
provider "google" {
  region = "us-central1"
}

# alternate configuration, whose alias is "europe"
provider "google" {
  alias  = "europe"
  region = "europe-west1"
}

resource "google_compute_instance" "example" {
  # This "provider" meta-argument selects the google provider
  # configuration whose alias is "europe", rather than the
  # default configuration.
  provider = google.europe

  # ...
}
```

You reference the alias under in the attribute **provider** for a resource.

# Introduction to Terraform Expressions

Expressions are used to **refer to** or **compute values** within a configuration.

Terraform Expressions is a large topic, and we'll be covering:

- Types and Values
- Strings and Templates
- References to Values
- Operators
- Function Calls
- Conditional Expressions
- For Expressions
- Splat Expressions
- Dynamic Blocks
- Type Constraints
- Version Constraints

# Types and Values

The result of an expression is a **value**. All values have a **type**

## primitive types

string

```
ami = "ami-830c94e3"
```

number

```
size = 6.283185
```

bool

```
termination_protection = true
```

## no type

null

```
endpoint = null
```

## complex/structural/collection types

list (tuple)

```
regions = ["us-east-1a", "us-east-1b"]
```

map (object)

```
tags = {env = "Production", priority = 3}
```

null represents *absence* or *omission*

*when you want to use the **underlying default** of a provider's resource configuration option*

# Strings

When quoting strings you use **double quotes** eg.

“hello”

Double quoted strings can interpret **escape sequences**.

|             |                                                     |
|-------------|-----------------------------------------------------|
| \n          | Newline                                             |
| \r          | Carriage Return                                     |
| \t          | Tab                                                 |
| \\"         | Literal quote (without terminating the string)      |
| \\\\"       | Literal backslash                                   |
| \uNNNN      | Unicode character from the basic multilingual plane |
| \UNNNNNNNNN | Unicode character from supplementary planes         |

Terraform also supports a "heredoc" style.  
Heredoc is a UNIX style multi-line string:

```
<<EOT  
hello  
world  
EOT
```

**special escape sequences:**

|        |                                                               |
|--------|---------------------------------------------------------------|
| \$\$\{ | Literal \${, without beginning an interpolation sequence.     |
| %%\{   | Literal %{}, without beginning a template directive sequence. |

# Strings Templates

String **interpolation** allows you to evaluate an expression between the markers eg.  `${....}` and converts it to a string.

"Hello, \${var.name}!"

String **directive** allows you to evaluate an conditional logic between the markers eg.  `%{....}`

"Hello, %{ if var.name != "" }\${var.name} %{ else }unnamed%{ endif }!"

You can use interpolation or directives within a **HEREDOC**

```
<<EOT
%{ for ip in aws_instance.example.*.private_ip }
server ${ip}
%{ endfor }
EOT
```

```
<<EOT
%{ for ip in aws_instance.example.*.private_ip ~}
server ${ip}
%{ endfor ~}
EOT
```

You can **strip whitespace** that would normally be left by directive blocks by providing a trailing tilde eg. `~`

# Expressions – Operators

Operators are **mathematical operations** you can preform to numbers within expressions

- Multiplication                    $a * b$
- Division                          $a / b$
- Modulus                          $a \% b$
- Addition                         $a + b$
- Subtraction                     $a - b$
- Flip to Negative (\* -1)     $-a$
- Equals                          $a == b$
- Does not Equal               $a != b$
- Less Than                      $a < b$
- Less Then or Equal         $a <= b$
- Greater Then                 $a > b$
- Greater Then or Equal     $a >= b$
- Or                              $a || b$
- And                             $a \&& b$
- Flip Boolean                 $!a$

# Conditional Expressions

Terraform support ternary if else conditions.

```
condition ? true_val : false_val
```

```
var.a != "" ? var.a : "default-a"
```

The return type for it the if and else must be the same type:

```
var.example ? tostring(12) : "hello"
```

# Expressions – For Expressions

For expressions allows you to iterate over a complex type and apply transformations

A for expression can accept as input **a list, a set, a tuple, a map, or an object**.

Uppercase each string in the provided list

```
[for s in var.list : upper(s)]
```

For map you can get: **Key** and **value**

```
[for k, v in var.map : length(k) + length(v)]
```

For a list you can get the **index**

```
[for i, v in var.list : "${i} is ${v}"]
```

Square braces **[ ]** returns a tuple

```
[for s in var.list : upper(s)]
```

→ [“HELLO”, WORLD”]

Curly braces **{ }** returns an object

```
{for s in var.list : s => upper(s)}
```

→ { hello = “HELLO”, world = “WORLD” }

# Expressions – For Expressions

An if statement can be used in a for expression to filter / reduce the amount of elements returned.

```
[for s in var.list : upper(s) if s != ""]
```

## Implicit Element Ordering on Conversion

Since Terraform can convert a unordered type (maps objects and sets) to a ordered type (list and tuples) it will need to choose an implied ordering.

- **Maps and Objects** – stored by key A-Z
- **Sets of Strings** – stored by string A-Z
- **Everything else** – arbitrary ordering

# Expressions – Splat Expressions

A **splat** expression provides a **shorter expression** for **for expressions**

**What is a splat operator?**

A splat operator is represented by an asterisk (\*), it originates from the ruby language

Splats in Terraform are used to rollup or soak up a bunch of iterations in a *for expression*

For lists, sets and tuples

```
[for o in var.list : o.id]
[for o in var.list : o.interfaces[0].name]
```



Can be written like this

```
var.list[*].id
var.list[*].interfaces[0].name
```

# Expressions – Splat Expressions

Splat expressions have a special behavior when you apply them to **lists**

- If the value is anything other than a null value then the splat expression will transform it into a single-element list
- If the value is *null* then the splat expression will return an empty tuple.

This behavior is useful for modules that accept optional input variables whose default value is *null* to represent the absence of any value to adapt the variable value to work with other Terraform language features that are designed to work with collections.

```
variable "website" {
  type = object({
    index_document = string
    error_document = string
  })
  default = null
}

resource "aws_s3_bucket" "example" {
  # ...

  dynamic "website" {
    for_each = var.website[*]
    content {
      index_document = website.value.index_document
      error_document = website.value.error_document
    }
  }
}
```

# Dynamic Blocks

Dynamic blocks allows you **dynamically construct repeatable nested blocks**

Lets say you need to create a bunch of ingress rules for an EC2 Security Group.



Define **objects** in locals:

Set **dynamic** block and →  
utilized for\_each block

```
locals {
    ingress_rules = [
        {
            port      = 443
            description = "Port 443"
        },
        {
            port      = 80
            description = "Port 80"
        }
    ]
}

resource "aws_security_group" "main" {
    name      = "sg"
    vpc_id   = data.aws_vpc.main.id

    dynamic "ingress" {
        for_each = local.ingress_rules

        content {
            description = ingress.value.description
            from_port   = ingress.value.port
            to_port     = ingress.value.port
            protocol    = "tcp"
            cidr_blocks = ["0.0.0.0/0"]
        }
    }
}
```

# Version Constraints

Terraform utilizes Semantic Versioning for specifying Terraform, Providers and Modules versions

Semantic Versioning is **open-standard** on how to define versioning for software management e.g. **MAJOR.MINOR.PATCH**



2.0.0    2.0.0-rc.2    2.0.0-rc.1    1.0.0    1.0.0-beta

[semver.org](https://semver.org)

1. **MAJOR** version when you make incompatible API changes,
2. **MINOR** version when you add functionality in a backwards compatible manner, and
3. **PATCH** version when you make backwards compatible bug fixes.

Additional labels for pre-release and build metadata are available as extensions to the MAJOR.MINOR.PATCH format.

A **version constraint** is a string containing one or more conditions, separated by commas.

- = or no operator. Match exact version number e.g. “1.0.0”, “=1.0.0”
- != Excludes an exact version number e.g. “!=1.0.0”
- > >= < <= Compare against a specific version e.g. “>= 1.0.0”
- ~> Allow only the rightmost version (last number) to increment e.g. “~> 1.0.0”

# Progressive Versioning

Progressive Versioning is the practice of using the latest version to keep a proactive stance of security, modernity and development agility



Practicing Good Hygiene

By being up to date you are pushing left things you will need to fix to stay compatible. You will have to deal with smaller problems instead of dealing with a big problem later on

**Run Nightly Builds** of your golden images or terraform plan as a warning signal to budget the time to improve for outage.

A nightly build is an automated workflow that occurs at night when developers are asleep. If the build breaks because a change is required for the code, the developers will see this upon arrival in the morning and be able to budget accordingly.

# Terraform State

## What is State?

A particular condition of cloud resources at a specific time.

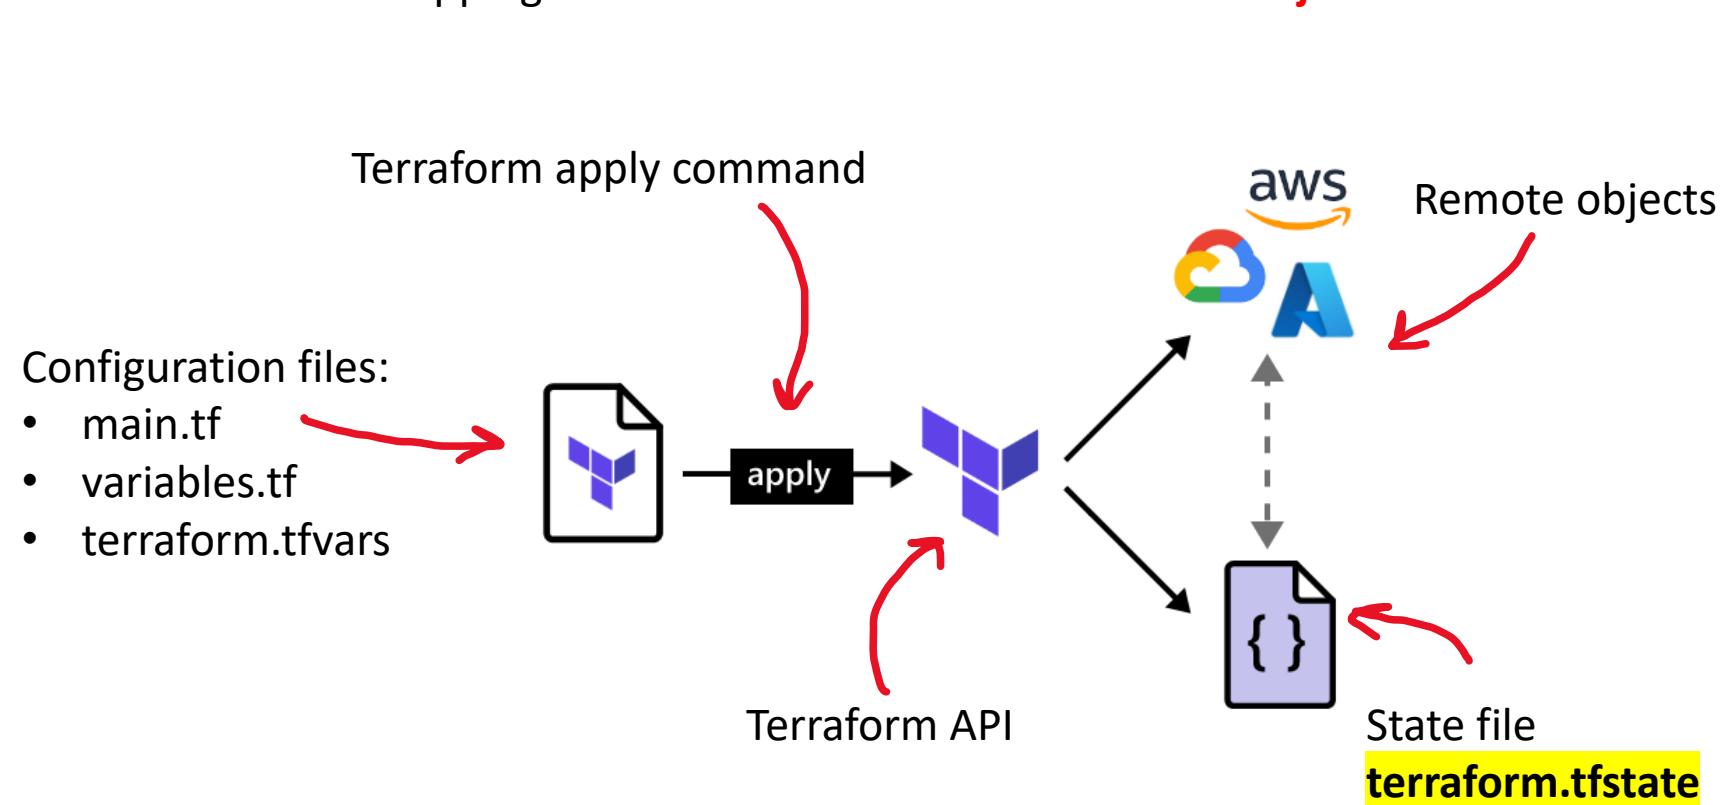
eg. We expect there to be a Virtual Machine running CentOS on AWS with a compute type of t2.micro.

## How does Terraform preserve state?

When you provision infrastructure via Terraform it will create a state file named **terraform.tfstate**

This **state file is a JSON data structure** with a one-to-one mapping from **resource instances** to **remote objects**

```
{  
  "version": 4,  
  "terraform_version": "1.0.4",  
  "serial": 1,  
  "lineage": "7b78e5ae-eae4-de2e-95a2-2a947f5d",  
  "outputs": {},  
  "resources": [  
    {  
      "mode": "managed",  
      "type": "aws_instance",  
      "name": "my_example_server",  
      "provider": "provider[\"registry.terraform.io/providers/hashicorp/aws/3.73.0\"]",  
      "instances": [  
        {  
          "schema_version": 1,  
          "attributes": {  
            "ami": "ami-0c2b8ca1dad447f8a",  
            "arn": "arn:aws:ec2:us-east-1:318479211178:instance/0d873544-1560-4311-b3e5-6a0a139885f5",  
            "associate_public_ip_address": true,  
            "availability_zone": "us-east-1a",  
            "block_device_mappings": [{"device_name": "/dev/sda1", "volume_id": "vol-0e108f53a3a44448"}, {"device_name": "nvme0n1", "volume_id": "vol-0a313b5f23a0441a"}],  
            "cpu_value": 0.25,  
            "ephemeral_block_devices": [{"device_name": "nvme0n1", "size_gb": 20}, {"device_name": "nvme1n1", "size_gb": 20}],  
            "iam_instance_profile": null,  
            "image_id": "ami-0c2b8ca1dad447f8a",  
            "instance_type": "t2.micro",  
            "key_name": null,  
            "lifecycle": null,  
            "load_balancers": null,  
            "name": "my-example-server",  
            "network_interface_ids": ["eni-0a313b5f23a0441a"],  
            "owner_id": "318479211178",  
            "placement_group": null,  
            "private_ip": "172.31.1.2",  
            "public_ip": "54.217.247.199",  
            "public_ipv4": "54.217.247.199",  
            "ramdisk_id": null,  
            "root_block_device": {"device_name": "/dev/sda1", "volume_id": "vol-0e108f53a3a4448"},  
            "root_device_name": "/dev/sda1",  
            "root_device_type": "partition",  
            "root_device_type": "partition",  
            "security_group_ids": ["sg-0b7a3f86a233a343"],  
            "source_dest_check": true,  
            "subnet_id": "subnet-0a313b5f23a0441a",  
            "tags": [{"key": "Name", "value": "my-example-server"}, {"key": "Environment", "value": "Production"}],  
            "vpc_security_group_ids": ["sg-0b7a3f86a233a343"]  
          }  
        ]  
      ]  
    }  
  ]  
}
```



# Terraform State

terraform state **list**

List resources in the state

terraform state **mv**

Move an item in the state

terraform state **pull**

Pull current remote state and output to stdout

terraform state **push**

Update remote state from a local state

terraform state **replace-provider**

Replace provider in the state

terraform state **rm**

Remove instances from the state

terraform state **show**

Show a resource in the state

# Terraform State Mv

**terraform state mv** allows you to:

- rename existing resources
- move a resource into a module
- move a module into a module

If you were to just rename a resource or move it to another module and run terraform apply Terraform will destroy and create the resource. State mv allows you to just change the reference so you can avoid a create and destroy action

Rename resource

```
terraform state mv packet_device.worker packet_device.helper
```

Move resource into module

```
terraform state mv packet_device.worker module.worker.packet_device.worker
```

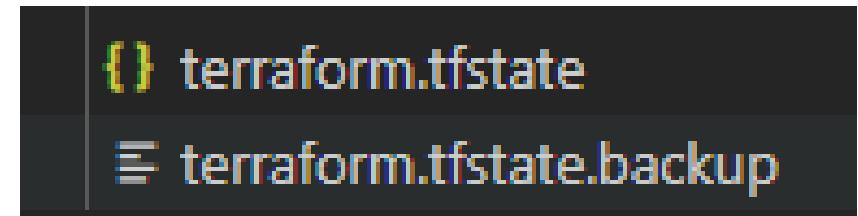
Move module into a module

```
terraform state mv module.app module.parent.module.app
```

# Terraform State Backups

All terraform state subcommands that *modify state* will write a backup file.  
Read only commands will not modify state eg. list, show

Terraform will take the current state and store  
it in a file called **terraform.tfstate.backup**



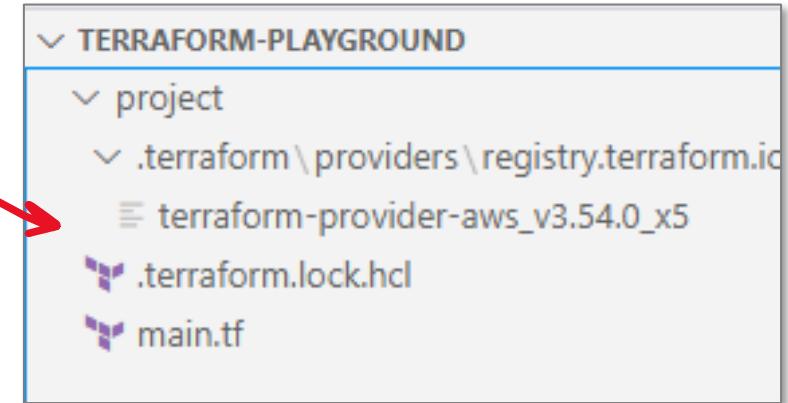
Backups cannot be disabled. This is by design to enforce best-practice for recovery

To get rid of the backup file you need to manually delete the files

# terraform Init

**terraform init** initializes your terraform project by:

- Downloading plugin dependencies e.g. Providers and Modules
- Create a .terraform directory
- Create a dependency lock file to enforce expected versions for plugins and terraform itself.



Terraform init is generally the first command you will run for a new terraform project  
If you modify or change dependencies run **terraform init** again to have it apply the changes

**terraform init -upgrade**

Upgrade all plugins to the latest version that  
complies with the configuration's version  
constraint  
Skip plugin installation

Dependency lock file  
**.terraform.lock.hcl**

**terraform init -get-plugins=false**

Force plugin installation to read plugins  
only from target directory

State lock file  
**.terraform.tfstate.lock.hcl**

**terraform init -plugin-dir=PATH**

Set a dependency lockfile mode

**terraform init -lockfile=MODE**

# terraform get

**terraform get** command is used to download and **update modules** in the root module.

When you're developer your own **Terraform Modules**

You may need to frequently pull updated modules but you do no want to initialize your state or pull new provider binaries.



Terraform Get is *lightweight* in this case because it only updates modules.

In most cases you want to use `terraform init`, with the exception of local module development

# Writing and Modifying Terraform Code

Terraform has three CLI commands that **improve debugging configuration scripts**:

`terraform fmt`

rewrites Terraform configuration files to a standard format and style

`terraform validate`

validates the syntax and arguments of the Terraform configuration files in a directory

`terraform console`

an interactive shell for evaluating Terraform expressions

# terraform fmt

This command applies a subset of the **Terraform language style conventions** along with other minor adjustments for readability

terraform fmt will by default look in the current directory and apply formatting to all .tf files.

adjusting spacing two spaces indent

```
provider "aws" {  
    profile = "exampro"  
    region  = "us-west-2"  
}
```



```
provider "aws" {  
    profile = "exampro"  
    region  = "us-west-2"  
}
```

syntax error

```
provider "aws"  
{  
    profile = "exampro"  
    region  = "us-west-2"  
}
```

terraform fmt --diff

```
main.tf  
--- old/main.tf  
+++ new/main.tf  
@@ -9,8 +9,8 @@  
 }
```

```
provider "aws" {  
-    profile = "exampro"  
-    region  = "us-west-2"  
+    profile = "exampro"  
+    region  = "us-west-2"  
}
```

Error: Invalid block definition

```
on main.tf line 11:  
11: provider "aws"  
12: {
```

A block definition must have block content delimited by "{" and "}", starting on the same line as the block header

# terraform validate

Terraform Validate runs checks that verify whether a configuration is syntactically valid and internally consistent, regardless of any provided variables or existing state

Validate is useful for general verification of reusable modules, including correctness of attribute names and value types

We are missing an `instance_type` which is a required when lauching an AWS Virtual Machine

```
$ terraform validate
Error: Missing required argument

with aws_instance.app_server,
on main.tf line 16, in resource "aws_instance" "app_server":
16: resource "aws_instance" "app_server" {

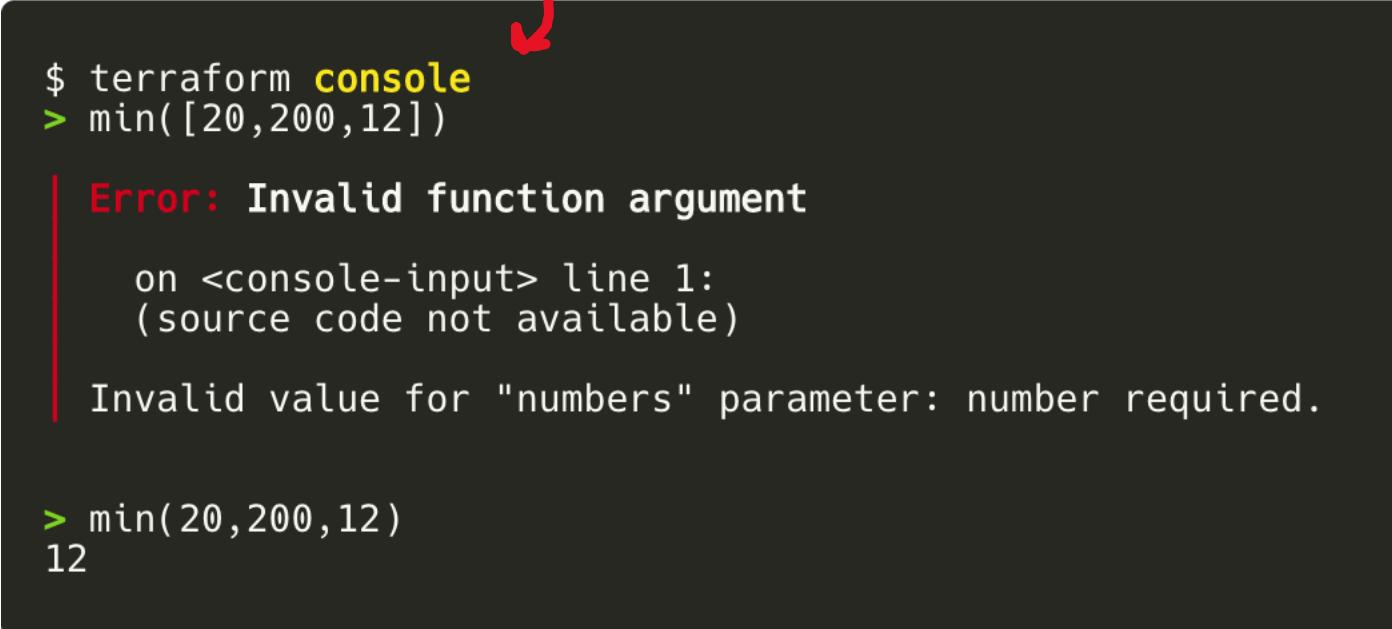
  "instance_type": one of `instance_type,launch_template` must be specified
```

When you run terraform **plan** or terraform **apply**, validate will automatically be performed.

# terraform console

Terraform console is an **interactive shell** where you can evaluate expressions.

We run a min command  
with the wrong arguments



A screenshot of the Terraform Console interface. The console has a dark background with white text. It shows the command '\$ terraform console' followed by the input 'min([20,200,12])'. A red arrow points from the text 'We run a min command with the wrong arguments' to the input line. Below the input, an error message is displayed in red: 'Error: Invalid function argument'. This is followed by two lines of explanatory text in white: 'on <console-input> line 1:' and '(source code not available)'. Another red arrow points from the text 'We correct the argument error' to the word 'numbers' in the error message. At the bottom of the console window, the corrected command is shown: '> min(20,200,12)' followed by the output '12'.

```
$ terraform console
> min([20,200,12])

Error: Invalid function argument

on <console-input> line 1:
(source code not available)

Invalid value for "numbers" parameter: number required.

> min(20,200,12)
12
```

We correct the  
argument error